

Flume 1.6.0 User Guide

Introduction

Overview

Apache Flume is a distributed, reliable, and available system for efficiently collecting, aggregating and moving large amounts of log data from many different sources to a centralized data store.

The use of Apache Flume is not only restricted to log data aggregation. Since data sources are customizable, Flume can be used to transport massive quantities of event data including but not limited to network traffic data, social-media-generated data, email messages and pretty much any data source possible.

Apache Flume is a top level project at the Apache Software Foundation.

There are currently two release code lines available, versions 0.9.x and 1.x.

Documentation for the 0.9.x track is available at [the Flume 0.9.x User Guide](#).

This documentation applies to the 1.4.x track.

New and existing users are encouraged to use the 1.x releases so as to leverage the performance improvements and configuration flexibilities available in the latest architecture.


System Requirements

1. Java Runtime Environment - Java 1.6 or later (Java 1.7 Recommended)
2. Memory - Sufficient memory for configurations used by sources, channels or sinks
3. Disk Space - Sufficient disk space for configurations used by channels or sinks
4. Directory Permissions - Read/Write permissions for directories used by agent

Architecture

Data flow model

A Flume event is defined as a unit of data flow having a byte payload and an optional set of string attributes. A Flume agent is a (JVM) process that hosts the components through which events flow from an external source to the next destination (hop).



A Flume source consumes events delivered to it by an external source like a web server. The external source sends events to Flume in a format that is recognized by the target Flume source. For example, an Avro Flume source can be used to receive Avro events from Avro clients or other Flume agents in the flow that send events from an Avro sink. A similar flow can be defined using a Thrift Flume Source to receive events from a Thrift Sink or a Flume Thrift Rpc Client or Thrift clients written in any language generated from the Flume thrift protocol. When a Flume source receives an event, it stores it into one or more channels. The channel is a passive store that keeps the event until it's consumed by a Flume sink. The file channel is one example – it is backed by the local filesystem. The sink removes the event from the channel and puts it into an external repository like HDFS (via Flume HDFS sink) or forwards it to the Flume source of the next Flume agent (next hop) in the flow. The source and sink within the given agent run asynchronously with the events staged in the channel.

Complex flows

Flume allows a user to build multi-hop flows where events travel through multiple agents before reaching the final destination. It also allows fan-in and fan-out flows, contextual routing and backup routes (fail-over) for failed hops.

Reliability

The events are staged in a channel on each agent. The events are then delivered to the next agent or terminal repository (like HDFS) in the flow. The events are removed from a channel only after

they are stored in the channel of next agent or in the terminal repository. This is how the single-hop message delivery semantics in Flume provide end-to-end reliability of the flow.

Flume uses a transactional approach to guarantee the reliable delivery of the events. The sources and sinks encapsulate in a transaction the storage/retrieval, respectively, of the events placed in or provided by a transaction provided by the channel. This ensures that the set of events are reliably passed from point to point in the flow. In the case of a multi-hop flow, the sink from the previous hop and the source from the next hop both have their transactions running to ensure that the data is safely stored in the channel of the next hop.

Recoverability

The events are staged in the channel, which manages recovery from failure. Flume supports a durable file channel which is backed by the local file system. There's also a memory channel which simply stores the events in an in-memory queue, which is faster but any events still left in the memory channel when an agent process dies can't be recovered.

Setup

Setting up an agent

Flume agent configuration is stored in a local configuration file. This is a text file that follows the Java properties file format. Configurations for one or more agents can be specified in the same configuration file. The configuration file includes properties of each source, sink and channel in an agent and how they are wired together to form data flows.

Configuring individual components

Each component (source, sink or channel) in the flow has a name, type, and set of properties that are specific to the type and instantiation. For example, an Avro source needs a hostname (or IP address) and a port number to receive data from. A memory channel can have max queue size ("capacity"), and an HDFS sink needs to know the file system URI, path to create files, frequency of file rotation ("hdfs.rollInterval") etc. All such attributes of a component needs to be set in the properties file of the hosting Flume agent.

Wiring the pieces together

The agent needs to know what individual components to load and how they are connected in order to constitute the flow. This is done by listing the names of each of the sources, sinks and channels in the agent, and then specifying the connecting channel for each sink and source. For example, an agent flows events from an Avro source called avroWeb to HDFS sink hdfs-cluster1 via a file channel called file-channel. The configuration file will contain names of these components and file-channel as a shared channel for both avroWeb source and hdfs-cluster1 sink.

Starting an agent

An agent is started using a shell script called flume-ng which is located in the bin directory of the Flume distribution. You need to specify the agent name, the config directory, and the config file on the command line:

```
$ bin/flume-ng agent -n $agent_name -c conf -f conf/flume-  
conf.properties.template
```

Now the agent will start running source and sinks configured in the given properties file.

A simple example

Here, we give an example configuration file, describing a single-node Flume deployment. This configuration lets a user generate events and subsequently logs them to the console.

```
# example.conf: A single-node Flume configuration  
  
# Name the components on this agent  
a1.sources = r1  
a1.sinks = k1  
a1.channels = c1  
  
# Describe/configure the source  
a1.sources.r1.type = netcat  
a1.sources.r1.bind = localhost  
a1.sources.r1.port = 44444  
  
# Describe the sink  
a1.sinks.k1.type = logger  
  
# Use a channel which buffers events in memory  
a1.channels.c1.type = memory  
a1.channels.c1.capacity = 1000  
a1.channels.c1.transactionCapacity = 100  
  
# Bind the source and sink to the channel
```

```
a1.sources.r1.channels = c1
a1.sinks.k1.channel = c1
```

This configuration defines a single agent named a1. a1 has a source that listens for data on port 44444, a channel that buffers event data in memory, and a sink that logs event data to the console. The configuration file names the various components, then describes their types and configuration parameters. A given configuration file might define several named agents; when a given Flume process is launched a flag is passed telling it which named agent to manifest.

Given this configuration file, we can start Flume as follows:

```
$ bin/flume-ng agent --conf conf --conf-file example.conf --name a1 -
Dflume.root.logger=INFO,console
```

Note that in a full deployment we would typically include one more option: `--conf=<conf-dir>`. The `<conf-dir>` directory would include a shell script *flume-env.sh* and potentially a log4j properties file. In this example, we pass a Java option to force Flume to log to the console and we go without a custom environment script.

From a separate terminal, we can then telnet port 44444 and send Flume an event:

```
$ telnet localhost 44444
Trying 127.0.0.1...
Connected to localhost.localdomain (127.0.0.1).
Escape character is '^]'.
Hello world! <ENTER>
OK
```

The original Flume terminal will output the event in a log message.

```
12/06/19 15:32:19 INFO source.NetcatSource: Source starting
12/06/19 15:32:19 INFO source.NetcatSource: Created
serverSocket:sun.nio.ch.ServerSocketChannelImpl[/127.0.0.1:44444]
12/06/19 15:32:34 INFO sink.LoggerSink: Event: { headers:{} body: 48 65 6C 6C
6F 20 77 6F 72 6C 64 21 0D      Hello world!. }
```

Congratulations - you've successfully configured and deployed a Flume agent! Subsequent sections cover agent configuration in much more detail.

Zookeeper based Configuration

Flume supports Agent configurations via Zookeeper. *This is an experimental feature.* The configuration file needs to be uploaded in the Zookeeper, under a configurable prefix. The configuration file is stored in Zookeeper Node data. Following is how the Zookeeper Node tree would look like for agents a1 and a2

```
- /flume
|- /a1 [Agent config file]
|- /a2 [Agent config file]
```

Once the configuration file is uploaded, start the agent with following options

```
$ bin/flume-ng agent -conf conf -z zkhost:2181,zkhost1:2181 -p /flume -name a1 -
Dflume.root.logger=INFO,console
```

| Argument Name | Default | Description |
|---------------|---------|--|
| z | – | Zookeeper connection string. Comma separated list of hostname:port |
| p | /flume | Base Path in Zookeeper to store Agent configurations |

Installing third-party plugins

Flume has a fully plugin-based architecture. While Flume ships with many out-of-the-box sources, channels, sinks, serializers, and the like, many implementations exist which ship separately from Flume.

While it has always been possible to include custom Flume components by adding their jars to the `FLUME_CLASSPATH` variable in the `flume-env.sh` file, Flume now supports a special directory called `plugins.d` which automatically picks up plugins that are packaged in a specific format. This allows for easier management of plugin packaging issues as well as simpler debugging and troubleshooting of several classes of issues, especially library dependency conflicts.

The plugins.d directory

The `plugins.d` directory is located at `$FLUME_HOME/plugins.d`. At startup time, the `flume-ng start` script looks in the `plugins.d` directory for plugins that conform to the below format and includes them in proper paths when starting up `java`.

Directory layout for plugins

Each plugin (subdirectory) within `plugins.d` can have up to three sub-directories:

1. lib - the plugin's jar(s)
2. libext - the plugin's dependency jar(s)
3. native - any required native libraries, such as .so files

Example of two plugins within the plugins.d directory:

```
plugins.d/  
plugins.d/custom-source-1/  
plugins.d/custom-source-1/lib/my-source.jar  
plugins.d/custom-source-1/libext/spring-core-2.5.6.jar  
plugins.d/custom-source-2/  
plugins.d/custom-source-2/lib/custom.jar  
plugins.d/custom-source-2/native/gettext.so
```

Data ingestion

Flume supports a number of mechanisms to ingest data from external sources.

RPC

An Avro client included in the Flume distribution can send a given file to Flume Avro source using avro RPC mechanism:

```
$ bin/flume-ng avro-client -H localhost -p 41414 -F /usr/logs/log.10
```

The above command will send the contents of /usr/logs/log.10 to to the Flume source listening on that ports.

Executing commands

There's an exec source that executes a given command and consumes the output. A single 'line' of output ie. text followed by carriage return ('\r') or line feed ('\n') or both together.

Note


Flume does not support tail as a source. One can wrap the tail command in an exec source to stream the file.

Network streams

Flume supports the following mechanisms to read data from popular log stream types, such as:

1. Avro
2. Thrift
3. Syslog
4. Netcat

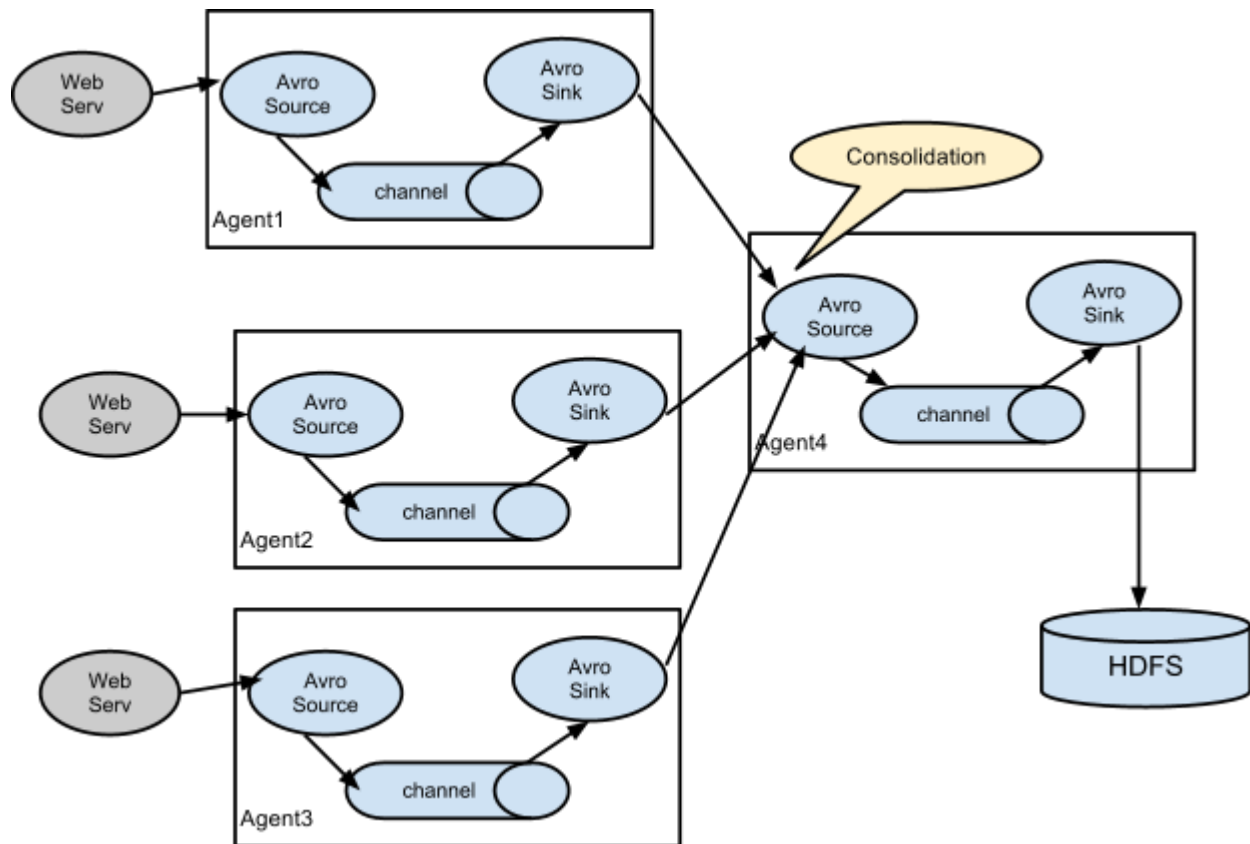
Setting multi-agent flow



In order to flow the data across multiple agents or hops, the sink of the previous agent and source of the current hop need to be avro type with the sink pointing to the hostname (or IP address) and port of the source.

Consolidation

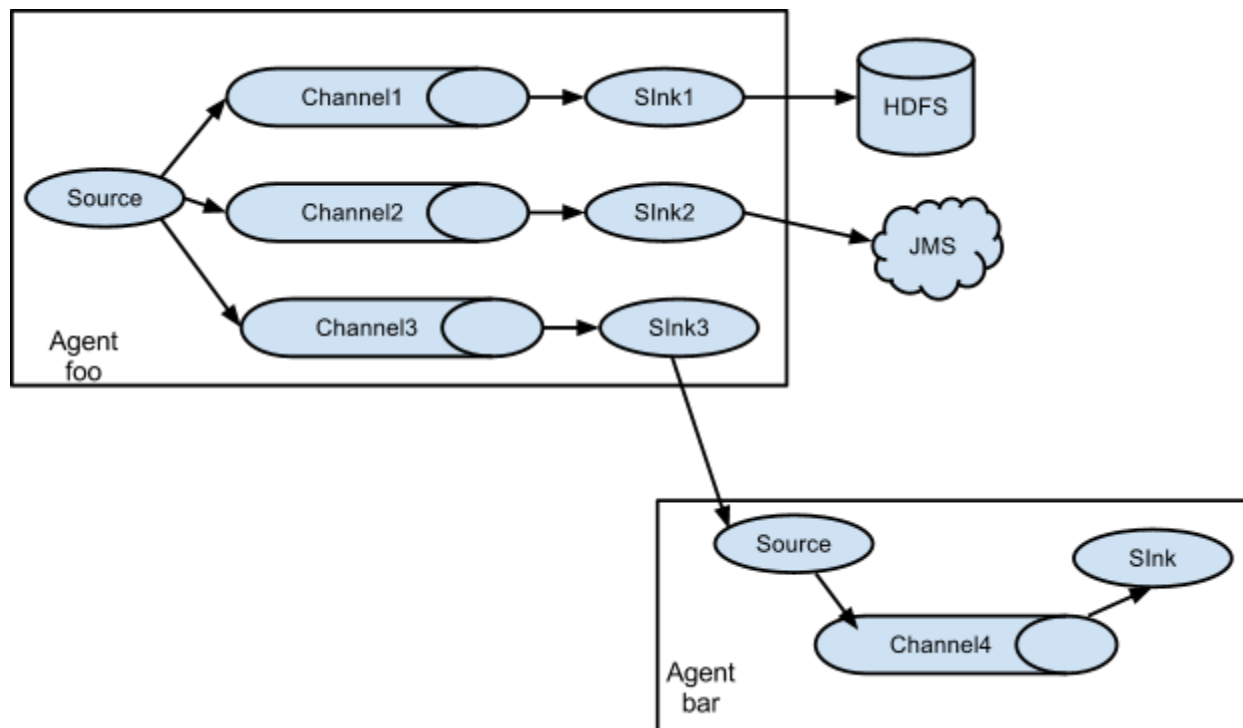
A very common scenario in log collection is a large number of log producing clients sending data to a few consumer agents that are attached to the storage subsystem. For example, logs collected from hundreds of web servers sent to a dozen of agents that write to HDFS cluster.



This can be achieved in Flume by configuring a number of first tier agents with an avro sink, all pointing to an avro source of single agent (Again you could use the thrift sources/sinks/clients in such a scenario). This source on the second tier agent consolidates the received events into a single channel which is consumed by a sink to its final destination.

Multiplexing the flow

Flume supports multiplexing the event flow to one or more destinations. This is achieved by defining a flow multiplexer that can replicate or selectively route an event to one or more channels.



The above example shows a source from agent “foo” fanning out the flow to three different channels. This fan out can be replicating or multiplexing. In case of replicating flow, each event is sent to all three channels. For the multiplexing case, an event is delivered to a subset of available channels when an event’s attribute matches a preconfigured value. For example, if an event attribute called “txnType” is set to “customer”, then it should go to channel1 and channel3, if it’s “vendor” then it should go to channel2, otherwise channel3. The mapping can be set in the agent’s configuration file.

Configuration

As mentioned in the earlier section, Flume agent configuration is read from a file that resembles a Java property file format with hierarchical property settings.

Defining the flow

To define the flow within a single agent, you need to link the sources and sinks via a channel. You need to list the sources, sinks and channels for the given agent, and then point the source and sink to a channel. A source instance can specify multiple channels, but a sink instance can only specify one channel. The format is as follows:

```
# list the sources, sinks and channels for the agent
<Agent>.sources = <Source>
<Agent>.sinks = <Sink>
<Agent>.channels = <Channel1> <Channel2>

# set channel for source
<Agent>.sources.<Source>.channels = <Channel1> <Channel2> ...

# set channel for sink
<Agent>.sinks.<Sink>.channel = <Channel1>
```

For example, an agent named `agent_foo` is reading data from an external avro client and sending it to HDFS via a memory channel. The config file `weblog.config` could look like:

```
# list the sources, sinks and channels for the agent
agent_foo.sources = avro-appserver-src-1
agent_foo.sinks = hdfs-sink-1
agent_foo.channels = mem-channel-1

# set channel for source
agent_foo.sources.avro-appserver-src-1.channels = mem-channel-1

# set channel for sink
agent_foo.sinks.hdfs-sink-1.channel = mem-channel-1
```

This will make the events flow from `avro-AppSrv-source` to `hdfs-Cluster1-sink` through the memory channel `mem-channel-1`. When the agent is started with the `weblog.config` as its config file, it will instantiate that flow.

Configuring individual components

After defining the flow, you need to set properties of each source, sink and channel. This is done in the same hierarchical namespace fashion where you set the component type and other values for the properties specific to each component:

```
# properties for sources
<Agent>.sources.<Source>.<someProperty> = <someValue>

# properties for channels
<Agent>.channel.<Channel>.<someProperty> = <someValue>

# properties for sinks
<Agent>.sources.<Sink>.<someProperty> = <someValue>
```

The property “type” needs to be set for each component for Flume to understand what kind of object it needs to be. Each source, sink and channel type has its own set of properties required for it to function as intended. All those need to be set as needed. In the previous example, we have a flow from avro-AppSrv-source to hdfs-Cluster1-sink through the memory channel mem-channel-1. Here’s an example that shows configuration of each of those components:

```
agent_foo.sources = avro-AppSrv-source
agent_foo.sinks = hdfs-Cluster1-sink
agent_foo.channels = mem-channel-1

# set channel for sources, sinks

# properties of avro-AppSrv-source
agent_foo.sources.avro-AppSrv-source.type = avro
agent_foo.sources.avro-AppSrv-source.bind = localhost
agent_foo.sources.avro-AppSrv-source.port = 10000

# properties of mem-channel-1
agent_foo.channels.mem-channel-1.type = memory
agent_foo.channels.mem-channel-1.capacity = 1000
agent_foo.channels.mem-channel-1.transactionCapacity = 100

# properties of hdfs-Cluster1-sink
agent_foo.sinks.hdfs-Cluster1-sink.type = hdfs
agent_foo.sinks.hdfs-Cluster1-sink.hdfs.path = hdfs://namenode/flume/webdata

#...
```

Adding multiple flows in an agent

A single Flume agent can contain several independent flows. You can list multiple sources, sinks and channels in a config. These components can be linked to form multiple flows:

```
# list the sources, sinks and channels for the agent
<Agent>.sources = <Source1> <Source2>
<Agent>.sinks = <Sink1> <Sink2>
<Agent>.channels = <Channel1> <Channel2>
```

Then you can link the sources and sinks to their corresponding channels (for sources) of channel (for sinks) to setup two different flows. For example, if you need to setup two flows in an agent, one going from an external avro client to external HDFS and another from output of a tail to avro sink, then here’s a config to do that:

```
# list the sources, sinks and channels in the agent
```

```
agent_foo.sources = avro-AppSrv-source1 exec-tail-source2
agent_foo.sinks = hdfs-Cluster1-sink1 avro-forward-sink2
agent_foo.channels = mem-channel-1 file-channel-2

# flow #1 configuration
agent_foo.sources.avro-AppSrv-source1.channels = mem-channel-1
agent_foo.sinks.hdfs-Cluster1-sink1.channel = mem-channel-1

# flow #2 configuration
agent_foo.sources.exec-tail-source2.channels = file-channel-2
agent_foo.sinks.avro-forward-sink2.channel = file-channel-2
```

Configuring a multi agent flow

To setup a multi-tier flow, you need to have an avro/thrift sink of first hop pointing to avro/thrift source of the next hop. This will result in the first Flume agent forwarding events to the next Flume agent. For example, if you are periodically sending files (1 file per event) using avro client to a local Flume agent, then this local agent can forward it to another agent that has the mounted for storage.

Weblog agent config:

```
# list sources, sinks and channels in the agent
agent_foo.sources = avro-AppSrv-source
agent_foo.sinks = avro-forward-sink
agent_foo.channels = file-channel

# define the flow
agent_foo.sources.avro-AppSrv-source.channels = file-channel
agent_foo.sinks.avro-forward-sink.channel = file-channel

# avro sink properties
agent_foo.sources.avro-forward-sink.type = avro
agent_foo.sources.avro-forward-sink.hostname = 10.1.1.100
agent_foo.sources.avro-forward-sink.port = 10000

# configure other pieces
#...
```

HDFS agent config:

```
# list sources, sinks and channels in the agent
agent_foo.sources = avro-collection-source
agent_foo.sinks = hdfs-sink
agent_foo.channels = mem-channel
```

```
# define the flow
agent_foo.sources.avro-collection-source.channels = mem-channel
agent_foo.sinks.hdfs-sink.channel = mem-channel

# avro sink properties
agent_foo.sources.avro-collection-source.type = avro
agent_foo.sources.avro-collection-source.bind = 10.1.1.100
agent_foo.sources.avro-collection-source.port = 10000

# configure other pieces
#...
```

Here we link the avro-forward-sink from the weblog agent to the avro-collection-source of the hdfs agent. This will result in the events coming from the external appserver source eventually getting stored in HDFS.

Fan out flow

As discussed in previous section, Flume supports fanning out the flow from one source to multiple channels. There are two modes of fan out, replicating and multiplexing. In the replicating flow, the event is sent to all the configured channels. In case of multiplexing, the event is sent to only a subset of qualifying channels. To fan out the flow, one needs to specify a list of channels for a source and the policy for the fanning it out. This is done by adding a channel “selector” that can be replicating or multiplexing. Then further specify the selection rules if it’s a multiplexer. If you don’t specify a selector, then by default it’s replicating:

```
# List the sources, sinks and channels for the agent
<Agent>.sources = <Source1>
<Agent>.sinks = <Sink1> <Sink2>
<Agent>.channels = <Channel1> <Channel2>

# set list of channels for source (separated by space)
<Agent>.sources.<Source1>.channels = <Channel1> <Channel2>

# set channel for sinks
<Agent>.sinks.<Sink1>.channel = <Channel1>
<Agent>.sinks.<Sink2>.channel = <Channel2>

<Agent>.sources.<Source1>.selector.type = replicating
```

The multiplexing select has a further set of properties to bifurcate the flow. This requires specifying a mapping of an event attribute to a set for channel. The selector checks for each configured attribute in the event header. If it matches the specified value, then that event is sent to all the

channels mapped to that value. If there's no match, then the event is sent to set of channels configured as default:

```
# Mapping for multiplexing selector
<Agent>.sources.<Source1>.selector.type = multiplexing
<Agent>.sources.<Source1>.selector.header = <someHeader>
<Agent>.sources.<Source1>.selector.mapping.<Value1> = <Channel1>
<Agent>.sources.<Source1>.selector.mapping.<Value2> = <Channel1> <Channel2>
<Agent>.sources.<Source1>.selector.mapping.<Value3> = <Channel2>
#...

<Agent>.sources.<Source1>.selector.default = <Channel2>
```

The mapping allows overlapping the channels for each value.

The following example has a single flow that multiplexed to two paths. The agent named agent_foo has a single avro source and two channels linked to two sinks:

```
# list the sources, sinks and channels in the agent
agent_foo.sources = avro-AppSrv-source1
agent_foo.sinks = hdfs-Cluster1-sink1 avro-forward-sink2
agent_foo.channels = mem-channel-1 file-channel-2

# set channels for source
agent_foo.sources.avro-AppSrv-source1.channels = mem-channel-1 file-channel-2

# set channel for sinks
agent_foo.sinks.hdfs-Cluster1-sink1.channel = mem-channel-1
agent_foo.sinks.avro-forward-sink2.channel = file-channel-2

# channel selector configuration
agent_foo.sources.avro-AppSrv-source1.selector.type = multiplexing
agent_foo.sources.avro-AppSrv-source1.selector.header = State
agent_foo.sources.avro-AppSrv-source1.selector.mapping.CA = mem-channel-1
agent_foo.sources.avro-AppSrv-source1.selector.mapping.AZ = file-channel-2
agent_foo.sources.avro-AppSrv-source1.selector.mapping.NY = mem-channel-1
file-channel-2
agent_foo.sources.avro-AppSrv-source1.selector.default = mem-channel-1
```

The selector checks for a header called "State". If the value is "CA" then its sent to mem-channel-1, if its "AZ" then it goes to file-channel-2 or if its "NY" then both. If the "State" header is not set or doesn't match any of the three, then it goes to mem-channel-1 which is designated as 'default'.

The selector also supports optional channels. To specify optional channels for a header, the config parameter 'optional' is used in the following way:

```
# channel selector configuration
agent_foo.sources.avro-AppSrv-source1.selector.type = multiplexing
agent_foo.sources.avro-AppSrv-source1.selector.header = State
agent_foo.sources.avro-AppSrv-source1.selector.mapping.CA = mem-channel-1
agent_foo.sources.avro-AppSrv-source1.selector.mapping.AZ = file-channel-2
agent_foo.sources.avro-AppSrv-source1.selector.mapping.NY = mem-channel-1
file-channel-2
agent_foo.sources.avro-AppSrv-source1.selector.optional.CA = mem-channel-1
file-channel-2
agent_foo.sources.avro-AppSrv-source1.selector.mapping.AZ = file-channel-2
agent_foo.sources.avro-AppSrv-source1.selector.default = mem-channel-1
```

The selector will attempt to write to the required channels first and will fail the transaction if even one of these channels fails to consume the events. The transaction is reattempted on **all** of the channels. Once all required channels have consumed the events, then the selector will attempt to write to the optional channels. A failure by any of the optional channels to consume the event is simply ignored and not retried.

If there is an overlap between the optional channels and required channels for a specific header, the channel is considered to be required, and a failure in the channel will cause the entire set of required channels to be retried. For instance, in the above example, for the header “CA” mem-channel-1 is considered to be a required channel even though it is marked both as required and optional, and a failure to write to this channel will cause that event to be retried on **all** channels configured for the selector.

Note that if a header does not have any required channels, then the event will be written to the default channels and will be attempted to be written to the optional channels for that header. Specifying optional channels will still cause the event to be written to the default channels, if no required channels are specified. If no channels are designated as default and there are no required, the selector will attempt to write the events to the optional channels. Any failures are simply ignored in that case.

Flume Sources

Avro Source

Listens on Avro port and receives events from external Avro client streams. When paired with the built-in Avro Sink on another (previous hop) Flume agent, it can create tiered collection topologies. Required properties are in **bold**.

| Property Name | Default | Description |
|-------------------|---------|---|
| channels | – | |
| type | – | The component type name, needs to be <code>avro</code> |
| bind | – | hostname or IP address to listen on |
| port | – | Port # to bind to |
| threads | – | Maximum number of worker threads to spawn |
| selector.type | | |
| selector.* | | |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |
| compression-type | none | This can be “none” or “deflate”. The compression-type must match the compression-type of matching AvroSource |
| ssl | false | Set this to true to enable SSL encryption. You must also specify a “keystore” and a “keystore-password”. |
| keystore | – | This is the path to a Java keystore file. Required for SSL. |
| keystore-password | – | The password for the Java keystore. Required for SSL. |
| keystore-type | JKS | The type of the Java keystore. This can be “JKS” or “PKCS12”. |
| exclude-protocols | SSLv3 | Space-separated list of SSL/TLS protocols to exclude. SSLv3 will always be excluded in addition to the protocols specified. |
| ipFilter | false | Set this to true to enable ipFiltering for netty |
| ipFilterRules | – | Define N netty ipFilter pattern rules with this config. |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = avro
a1.sources.r1.channels = c1
a1.sources.r1.bind = 0.0.0.0
a1.sources.r1.port = 4141
```

Example of ipFilterRules

ipFilterRules defines N netty ipFilters separated by a comma a pattern rule must be in this format.

<'allow' or deny>:<'ip' or 'name' for computer name>:<pattern> or allow/deny:ip/name:pattern

example: ipFilterRules=allow:ip:127.*;allow:name:localhost;deny:ip:*

Note that the first rule to match will apply as the example below shows from a client on the localhost

This will Allow the client on localhost be deny clients from any other ip
“allow:name:localhost,deny:ip:” This will deny the client on localhost be allow clients from any other ip
“deny:name:localhost,allow:ip:”

Thrift Source

Listens on Thrift port and receives events from external Thrift client streams. When paired with the built-in ThriftSink on another (previous hop) Flume agent, it can create tiered collection topologies. Thrift source can be configured to start in secure mode by enabling kerberos authentication. agent-principal and agent-keytab are the properties used by the Thrift source to authenticate to the kerberos KDC. Required properties are in **bold**.

| Property Name | Default | Description |
|-------------------|---------|---|
| channels | – | |
| type | – | The component type name, needs to be <code>thrift</code> |
| bind | – | hostname or IP address to listen on |
| port | – | Port # to bind to |
| threads | – | Maximum number of worker threads to spawn |
| selector.type | | |
| selector.* | | |
| interceptors | – | Space separated list of interceptors |
| interceptors.* | | |
| ssl | false | Set this to true to enable SSL encryption. You must also specify a “keystore” and a “keystore-password”. |
| keystore | – | This is the path to a Java keystore file. Required for SSL. |
| keystore-password | – | The password for the Java keystore. Required for SSL. |
| keystore-type | JKS | The type of the Java keystore. This can be “JKS” or “PKCS12”. |
| exclude-protocols | SSLv3 | Space-separated list of SSL/TLS protocols to exclude. SSLv3 will always be excluded in addition to the protocols specified. |
| kerberos | false | Set to true to enable kerberos authentication. In kerberos mode, agent-principal and agent-keytab are required for successful authentication. The Thrift source in secure mode, will accept connections only from Thrift clients that have kerberos enabled and are successfully authenticated to the kerberos KDC. |

| Property Name | Default | Description |
|-----------------|---------|--|
| agent-principal | — | The kerberos principal used by the Thrift Source to authenticate to the kerberos KDC. |
| agent-keytab | — | The keytab location used by the Thrift Source in combination with the agent-principal to authenticate to the kerberos KDC. |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = thrift
a1.sources.r1.channels = c1
a1.sources.r1.bind = 0.0.0.0
a1.sources.r1.port = 4141
```

Exec Source

Exec source runs a given Unix command on start-up and expects that process to continuously produce data on standard out (stderr is simply discarded, unless property logStdErr is set to true). If the process exits for any reason, the source also exits and will produce no further data. This means configurations such as `cat [named pipe]` or `tail -F [file]` are going to produce the desired results where as `date` will probably not - the former two commands produce streams of data where as the latter produces a single event and exits.

Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------|---------|--|
| channels | — | |
| type | — | The component type name, needs to be <code>exec</code> |
| command | — | The command to execute |
| shell | — | A shell invocation used to run the command. e.g. <code>/bin/sh -c</code> . Required only for commands relying on shell features like wildcards, back ticks, pipes etc. |
| restartThrottle | 10000 | Amount of time (in millis) to wait before attempting a restart |
| restart | false | Whether the executed cmd should be restarted if it dies |
| logStdErr | false | Whether the command's stderr should be logged |
| batchSize | 20 | The max number of lines to read and send to the channel at a time |

| Property Name | Default | Description |
|----------------|-------------|--|
| batchTimeout | 3000 | Amount of time (in milliseconds) to wait, if the buffer size was not reached, before data is pushed downstream |
| selector.type | replicating | replicating or multiplexing |
| selector.* | | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |

Warning

The problem with ExecSource and other asynchronous sources is that the source can not guarantee that if there is a failure to put the event into the Channel the client knows about it. In such cases, the data will be lost. As a for instance, one of the most commonly requested features is the `tail -F [file]`-like use case where an application writes to a log file on disk and Flume tails the file, sending each line as an event. While this is possible, there's an obvious problem; what happens if the channel fills up and Flume can't send an event? Flume has no way of indicating to the application writing the log file that it needs to retain the log or that the event hasn't been sent, for some reason. If this doesn't make sense, you need only know this: Your application can never guarantee data has been received when using a unidirectional asynchronous interface such as ExecSource! As an extension of this warning - and to be completely clear - there is absolutely zero guarantee of event delivery when using this source. For stronger reliability guarantees, consider the Spooling Directory Source or direct integration with Flume via the SDK.

Note

You can use ExecSource to emulate TailSource from Flume 0.9x (flume og). Just use unix command `tail -F /full/path/to/your/file`. Parameter -F is better in this case than -f as it will also follow file rotation.

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = exec
a1.sources.r1.command = tail -F /var/log/secure
a1.sources.r1.channels = c1
```

The 'shell' config is used to invoke the 'command' through a command shell (such as Bash or Powershell). The 'command' is passed as an argument to 'shell' for execution. This allows the

‘command’ to use features from the shell such as wildcards, back ticks, pipes, loops, conditionals etc. In the absence of the ‘shell’ config, the ‘command’ will be invoked directly. Common values for ‘shell’ : ‘/bin/sh -c’, ‘/bin/ksh -c’, ‘cmd /c’, ‘powershell -Command’, etc.

```
al.sources.tailsource-1.type = exec
al.sources.tailsource-1.shell = /bin/bash -c
al.sources.tailsource-1.command = for i in /path/*.txt; do cat $i; done
```

JMS Source

JMS Source reads messages from a JMS destination such as a queue or topic. Being a JMS application it should work with any JMS provider but has only been tested with ActiveMQ. The JMS source provides configurable batch size, message selector, user/pass, and message to flume event converter. Note that the vendor provided JMS jars should be included in the Flume classpath using plugins.d directory (preferred), –classpath on command line, or via FLUME_CLASSPATH variable in flume-env.sh.

Required properties are in **bold**.

| Property Name | Default | Description |
|------------------------------|---------|--|
| channels | – | |
| type | – | The component type name, needs to be <code>jms</code> |
| initialContextFactory | – | Initial Context Factory, e.g: <code>org.apache.activemq.jndi.ActiveMQInitialContextFactory</code> |
| connectionFactory | – | The JNDI name the connection factory should appear as |
| providerURL | – | The JMS provider URL |
| destinationName | – | Destination name |
| destinationType | – | Destination type (queue or topic) |
| messageSelector | – | Message selector to use when creating the consumer |
| userName | – | Username for the destination/provider |
| passwordFile | – | File containing the password for the destination/provider |
| batchSize | 100 | Number of messages to consume in one batch |
| converter.type | DEFAULT | Class to use to convert messages to flume events. See below. |
| converter.* | – | Converter properties. |
| converter.charset | UTF-8 | Default converter only. Charset to use when converting JMS TextMessages to byte arrays. |

Converter

The JMS source allows pluggable converters, though it's likely the default converter will work for most purposes. The default converter is able to convert Bytes, Text, and Object messages to FlumeEvents. In all cases, the properties in the message are added as headers to the FlumeEvent.

BytesMessage:

Bytes of message are copied to body of the FlumeEvent. Cannot convert more than 2GB of data per message.

TextMessage:

Text of message is converted to a byte array and copied to the body of the FlumeEvent. The default converter uses UTF-8 by default but this is configurable.

ObjectMessage:

Object is written out to a ByteArrayOutputStream wrapped in an ObjectOutputStream and the resulting array is copied to the body of the FlumeEvent.

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = jms
a1.sources.r1.channels = c1
a1.sources.r1.initialContextFactory =
org.apache.activemq.jndi.ActiveMQInitialContextFactory
a1.sources.r1.connectionFactory = GenericConnectionFactory
a1.sources.r1.providerURL = tcp://mqserver:61616
a1.sources.r1.destinationName = BUSINESS_DATA
a1.sources.r1.destinationType = QUEUE
```

Spooling Directory Source

This source lets you ingest data by placing files to be ingested into a “spooling” directory on disk. This source will watch the specified directory for new files, and will parse events out of new files as they appear. The event parsing logic is pluggable. After a given file has been fully read into the channel, it is renamed to indicate completion (or optionally deleted).

Unlike the Exec source, this source is reliable and will not miss data, even if Flume is restarted or killed. In exchange for this reliability, only immutable, uniquely-named files must be dropped into the spooling directory. Flume tries to detect these problem conditions and will fail loudly if they are violated:

1. If a file is written to after being placed into the spooling directory, Flume will print an error to its log file and stop processing.
2. If a file name is reused at a later time, Flume will print an error to its log file and stop processing.

To avoid the above issues, it may be useful to add a unique identifier (such as a timestamp) to log file names when they are moved into the spooling directory.

Despite the reliability guarantees of this source, there are still cases in which events may be duplicated if certain downstream failures occur. This is consistent with the guarantees offered by other Flume components.

| Property Name | Default | Description |
|--------------------------|--------------------------|--|
| channels | – | |
| type | – | The component type name, needs to be <code>spoolDir</code> . |
| spoolDir | – | The directory from which to read files from. |
| fileSuffix | <code>.COMPLETE.D</code> | Suffix to append to completely ingested files |
| deletePolicy | <code>never</code> | When to delete completed files: <code>never</code> or <code>immediate</code> |
| fileHeader | <code>false</code> | Whether to add a header storing the absolute path filename. |
| fileHeaderKey | <code>file</code> | Header key to use when appending absolute path filename to event header. |
| basenameHeader | <code>false</code> | Whether to add a header storing the basename of the file. |
| basenameHeaderKey | <code>basename</code> | Header Key to use when appending basename of file to event header. |
| ignorePattern | <code>^\$</code> | Regular expression specifying which files to ignore (skip) |
| trackerDir | <code>.flumespool</code> | Directory to store metadata related to processing of files. If this path is not an absolute path, then it is interpreted as relative to the <code>spoolDir</code> . |
| consumeOrder | <code>oldest</code> | In which order files in the spooling directory will be consumed <code>oldest</code> , <code>youngest</code> and <code>random</code> . In case of <code>oldest</code> and <code>youngest</code> , the last modified time of the files will be |

| Property Name | Default | Description |
|----------------------------------|---------|---|
| | | used to compare the files. In case of a tie, the file with smallest lxicographical order will be consumed first. In case of <code>random</code> any file will be picked randomly. When using <code>oldest</code> and <code>youngest</code> the whole directory will be scanned to pick the oldest/youngest file, which might be slow if there are a large number of files, while using <code>random</code> may cause old files to be consumed very late if new files keep coming in the spooling directory. |
| <code>maxBackoff</code> | 4000 | The maximum time (in millis) to wait between consecutive attempts to write to the channel(s) if the channel is full. The source will start at a low backoff and increase it exponentially each time the channel throws a <code>ChannelException</code> , upto the value specified by this parameter. |
| <code>batchSize</code> | 100 | Granularity at which to batch transfer to the channel |
| <code>inputCharset</code> | UTF-8 | Character set used by deserializers that treat the input file as text. |
| <code>decodeErrorPolicy</code> | FAIL | What to do when we see a non-decodable character in the input file. <code>FAIL</code> : Throw an exception and fail to parse the file. <code>REPLACE</code> : Replace the unparseable character with the “replacement character” char, typically Unicode <code>U+FFFD</code> . <code>IGNORE</code> : Drop the unparseable character sequence. |
| <code>deserializer</code> | LINE | Specify the deserializer used to parse the file into events. Defaults to parsing each line as an event. The class specified must implement <code>EventDeserializer.Builder</code> . |
| <code>deserializer.*</code> | | Varies per event deserializer. |
| <code>bufferMaxLines</code> | – | (Obselete) This option is now ignored. |
| <code>bufferMaxLineLength</code> | 5000 | (Deprecated) Maximum length of a line in the commit buffer. Use <code>deserializer.maxLineLength</code> instead. |

| Property Name | Default | Description |
|----------------|-------------|--------------------------------------|
| selector.type | replicating | replicating or multiplexing |
| selector.* | | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |

Example for an agent named agent-1:

```
a1.channels = ch-1
a1.sources = src-1

a1.sources.src-1.type = spooldir
a1.sources.src-1.channels = ch-1
a1.sources.src-1.spoolDir = /var/log/apache/flumeSpool
a1.sources.src-1.fileHeader = true
```

Event Deserializers

The following event deserializers ship with Flume.

LINE

This deserializer generates one event per line of text input.

| Property Name | Default | Description |
|----------------------------|---------|--|
| deserializer.maxLineLength | 2048 | Maximum number of characters to include in a single event. If a line exceeds this length, it is truncated, and the remaining characters on the line will appear in a subsequent event. |
| deserializer.outputCharset | UTF-8 | Charset to use for encoding events put into the channel. |

AVRO

This deserializer is able to read an Avro container file, and it generates one event per Avro record in the file. Each event is annotated with a header that indicates the schema used. The body of the event is the binary Avro record data, not including the schema or the rest of the container file elements.

Note that if the spool directory source must retry putting one of these events onto a channel (for example, because the channel is full), then it will reset and retry from the most recent Avro container file sync point. To reduce potential event duplication in such a failure scenario, write sync markers more frequently in your Avro input files.

| Property Name | Default | Description |
|-------------------------|---------|---|
| deserializer.schemaType | HASH | How the schema is represented. By default, or when the value <code>HASH</code> is specified, the Avro schema is hashed and the hash is stored in every event in the event header “flume.avro.schema.hash”. If <code>LITERAL</code> is specified, the JSON-encoded schema itself is stored in every event in the event header “flume.avro.schema.literal”. Using <code>LITERAL</code> mode is relatively inefficient compared to <code>HASH</code> mode. |

BlobDeserializer

This deserializer reads a Binary Large Object (BLOB) per event, typically one BLOB per file. For example a PDF or JPG file. Note that this approach is not suitable for very large objects because the entire BLOB is buffered in RAM.

| Property Name | Default | Description |
|----------------------------|----------|---|
| deserializer | – | The FQCN of this class: <code>org.apache.flume.sink.solr.morphline.BlobDeserializer\$Builder</code> |
| deserializer.maxBlobLength | 10000000 | The maximum number of bytes to read and buffer for a given request |

Twitter 1% firehose Source (experimental)

Warning

This source is highly experimental and may change between minor versions of Flume. Use at your own risk.

Experimental source that connects via Streaming API to the 1% sample twitter firehose, continuously downloads tweets, converts them to Avro format and sends Avro events to a downstream Flume sink. Requires the consumer and access tokens and secrets of a Twitter developer account. Required properties are in **bold**.

| Property Name | Default | Description |
|-------------------------------|---------|---|
| channels | – | |
| type | – | The component type name, needs to be <code>org.apache.flume.source.twitter.TwitterSource</code> |
| consumerKey | – | OAuth consumer key |
| consumerSecret | – | OAuth consumer secret |
| accessToken | – | OAuth access token |
| accessTokenSecret | – | OAuth token secret |
| maxBatchSize | 1000 | Maximum number of twitter messages to put in a single batch |
| maxBatchDurationMillis | 1000 | Maximum number of milliseconds to wait before closing a batch |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = org.apache.flume.source.twitter.TwitterSource
a1.sources.r1.channels = c1
a1.sources.r1.consumerKey = YOUR_TWITTER_CONSUMER_KEY
a1.sources.r1.consumerSecret = YOUR_TWITTER_CONSUMER_SECRET
a1.sources.r1.accessToken = YOUR_TWITTER_ACCESS_TOKEN
a1.sources.r1.accessTokenSecret = YOUR_TWITTER_ACCESS_TOKEN_SECRET
a1.sources.r1.maxBatchSize = 10
a1.sources.r1.maxBatchDurationMillis = 200
```

Kafka Source

Kafka Source is an Apache Kafka consumer that reads messages from a Kafka topic. If you have multiple Kafka sources running, you can configure them with the same Consumer Group so each will read a unique set of partitions for the topic.

| Property Name | Default | Description |
|--|---------|---|
| channels | – | |
| type | – | The component type name, needs to be <code>org.apache.flume.source.kafka,KafkaSource</code> |
| zookeeperConnect | – | URI of ZooKeeper used by Kafka cluster |
| groupId | flume | Unique identified of consumer group. Setting the same id in multiple sources or agents indicates that they are part of the same consumer group |
| topic | – | Kafka topic we'll read messages from. At the time, this is a single topic only. |
| batchSize | 1000 | Maximum number of messages written to Channel in one batch |
| batchDurationMillis | 1000 | Maximum time (in ms) before a batch will be written to Channel The batch will be written whenever the first of size and time will be reached. |
| backoffSleepIncrement | 1000 | Initial and incremental wait time that is triggered when a Kafka Topic appears to be empty. Wait period will reduce aggressive pinging of an empty Kafka Topic. One second is ideal for ingestion use cases but a lower value may be required for low latency operations with interceptors. |
| maxBackoffSleep | 5000 | Maximum wait time that is triggered when a Kafka Topic appears to be empty. Five seconds is ideal for ingestion use cases but a lower value may be required for low latency operations with interceptors. |
| Other Kafka Consumer Properties | – | These properties are used to configure the Kafka Consumer. Any producer property supported by Kafka can be used. The only requirement is to prepend the property name with the prefix <code>kafka..</code> . For example: <code>kafka.consumer.timeout.ms</code> Check <i>Kafka documentation</i> < https://kafka.apache.org/08/configuration.html#consumerconfigs > for details |

Note

The Kafka Source overrides two Kafka consumer parameters: `auto.commit.enable` is set to “false” by the source and we commit every batch. For improved performance this can be set to “true”, however, this can lead to loss of data `consumer.timeout.ms`

is set to 10ms, so when we check Kafka for new data we wait at most 10ms for the data to arrive setting this to a higher value can reduce CPU utilization (we'll poll Kafka in less of a tight loop), but also means higher latency in writing batches to channel (since we'll wait longer for data to arrive).

Example for agent named tier1:

```
tier1.sources.source1.type =  
org.apache.flume.source.kafka.KafkaSource  
tier1.sources.source1.channels = channel1  
tier1.sources.source1.zookeeperConnect = localhost:2181  
tier1.sources.source1.topic = test1  
tier1.sources.source1.groupId = flume  
tier1.sources.source1.kafka.consumer.timeout.ms = 100
```

NetCat Source

A netcat-like source that listens on a given port and turns each line of text into an event. Acts like `nc -k -l [host] [port]`. In other words, it opens a specified port and listens for data. The expectation is that the supplied data is newline separated text. Each line of text is turned into a Flume event and sent via the connected channel.

Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------|-------------|--|
| channels | — | |
| type | — | The component type name, needs to be <code>netcat</code> |
| bind | — | Host name or IP address to bind to |
| port | — | Port # to bind to |
| max-line-length | 512 | Max line length per event body (in bytes) |
| ack-every-event | true | Respond with an "OK" for every event received |
| selector.type | replicating | replicating or multiplexing |
| selector.* | | Depends on the selector.type value |
| interceptors | — | Space-separated list of interceptors |
| interceptors.* | | |

Example for agent named a1:

```
a1.sources = r1  
a1.channels = c1
```

```
a1.sources.r1.type = netcat  
a1.sources.r1.bind = 0.0.0.0  
a1.sources.r1.bind = 6666  
a1.sources.r1.channels = c1
```

Sequence Generator Source

A simple sequence generator that continuously generates events with a counter that starts from 0 and increments by 1. Useful mainly for testing. Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------|-------------|--|
| channels | – | |
| type | – | The component type name, needs to be seq |
| selector.type | | replicating or multiplexing |
| selector.* | replicating | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |
| batchSize | 1 | |

Example for agent named a1:

```
a1.sources = r1  
a1.channels = c1  
a1.sources.r1.type = seq  
a1.sources.r1.channels = c1
```

Syslog Sources

Reads syslog data and generate Flume events. The UDP source treats an entire message as a single event. The TCP sources create a new event for each string of characters separated by a newline ('n').

Required properties are in **bold**.

Syslog TCP Source

The original, tried-and-true syslog TCP source.

| Property Name | Default | Description |
|-----------------------|-------------|---|
| channels | – | |
| type | – | The component type name, needs to be <code>syslogtcp</code> |
| host | – | Host name or IP address to bind to |
| port | – | Port # to bind to |
| eventSize | 2500 | Maximum size of a single event line, in bytes |
| keepFields | none | Setting this to 'all' will preserve the Priority, Timestamp and Hostname in the body of the event. A spaced separated list of fields to include is allowed as well. Currently, the following fields can be included: priority, version, timestamp, hostname. The values 'true' and 'false' have been deprecated in favor of 'all' and 'none'. |
| selector.type | | replicating or multiplexing |
| selector.* | replicating | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |

For example, a syslog TCP source for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = syslogtcp
a1.sources.r1.port = 5140
a1.sources.r1.host = localhost
a1.sources.r1.channels = c1
```

Multiport Syslog TCP Source

This is a newer, faster, multi-port capable version of the Syslog TCP source. Note that the `ports` configuration setting has replaced `port`. Multi-port capability means that it can listen on many ports at once in an efficient manner. This source uses the Apache Mina library to do that. Provides support for RFC-3164 and many common RFC-5424 formatted messages. Also provides the capability to configure the character set used on a per-port basis.

| Property Name | Default | Description |
|-----------------|---------|-------------|
| channels | – | |

| Property Name | Default | Description |
|----------------------------------|-----------------|---|
| type | – | The component type name, needs to be <code>multiport_syslogtcp</code> |
| host | – | Host name or IP address to bind to. |
| ports | – | Space-separated list (one or more) of ports to bind to. |
| eventSize | 2500 | Maximum size of a single event line, in bytes. |
| keepFields | none | Setting this to 'all' will preserve the Priority, Timestamp and Hostname in the body of the event. A spaced separated list of fields to include is allowed as well. Currently, the following fields can be included: priority, version, timestamp, hostname. The values 'true' and 'false' have been deprecated in favor of 'all' and 'none'. |
| portHeader | – | If specified, the port number will be stored in the header of each event using the header name specified here. This allows for interceptors and channel selectors to customize routing logic based on the incoming port. |
| charset.default | UTF-8 | Default character set used while parsing syslog events into strings. |
| charset.port.<port> | – | Character set is configurable on a per-port basis. |
| batchSize | 100 | Maximum number of events to attempt to process per request loop. Using the default is usually fine. |
| readBufferSize | 1024 | Size of the internal Mina read buffer. Provided for performance tuning. Using the default is usually fine. |
| numProcessors | (auto-detected) | Number of processors available on the system for use while processing messages. Default is to auto-detect # of CPUs using the Java Runtime API. Mina will spawn 2 request-processing threads per detected CPU, which is often reasonable. |
| selector.type | replicating | replicating, multiplexing, or custom |
| selector.* | – | Depends on the <code>selector.type</code> value |
| interceptors | – | Space-separated list of interceptors. |
| interceptors.* | | |

For example, a multiport syslog TCP source for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = multiport_syslogtcp
a1.sources.r1.channels = c1
a1.sources.r1.host = 0.0.0.0
a1.sources.r1.ports = 10001 10002 10003
a1.sources.r1.portHeader = port
```

Syslog UDP Source

| Property Name | Default | Description |
|----------------|-------------|---|
| channels | – | |
| type | – | The component type name, needs to be syslogudp |
| host | – | Host name or IP address to bind to |
| port | – | Port # to bind to |
| keepFields | false | Setting this to true will preserve the Priority, Timestamp and Hostname in the body of the event. |
| selector.type | | replicating or multiplexing |
| selector.* | replicating | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |

For example, a syslog UDP source for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = syslogudp
a1.sources.r1.port = 5140
a1.sources.r1.host = localhost
a1.sources.r1.channels = c1
```

HTTP Source

A source which accepts Flume Events by HTTP POST and GET. GET should be used for experimentation only. HTTP requests are converted into flume events by a pluggable “handler” which must implement the HTTPSourceHandler interface. This handler takes a HttpServletRequest and returns a list of flume events. All events handled from one Http request are committed to the channel in one transaction, thus

allowing for increased efficiency on channels like the file channel. If the handler throws an exception, this source will return a HTTP status of 400. If the channel is full, or the source is unable to append events to the channel, the source will return a HTTP 503 - Temporarily unavailable status.

All events sent in one post request are considered to be one batch and inserted into the channel in one transaction.

| Property Name | Default | Description |
|---------------|--|---|
| type | | The component type name, needs to be http |
| port | – | The port the source should bind to. |
| bind | 0.0.0.0 | The hostname or IP address to listen on |
| handler | org.apache.flume.source.http.JSONHandler | The FQCN of the handler class. |
| handler.* | – | Config parameters for the handler |
| selector.type | replicating | replicating or multiplexing |
| selector.* | | Depends on the selector.type value |
| interceptors | – | Space-separated |

| Property Name | Default | Description |
|------------------|-------------------|--|
| | | list of interceptors |
| interceptors.* | | |
| enableSSL | false | Set the property true, to enable SSL. <i>HTTP Source does not support SSLv3.</i> |
| excludeProtocols | SSLv3 | Space-separated list of SSL/TLS protocols to exclude. SSLv3 is always excluded. |
| keystore | | Location of the keystore including keystore file name |
| keystorePassword | Keystore password | |

For example, a http source for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = http
a1.sources.r1.port = 5140
a1.sources.r1.channels = c1
a1.sources.r1.handler = org.example.rest.RestHandler
a1.sources.r1.handler.nickname = random props
```

JSONHandler

A handler is provided out of the box which can handle events represented in JSON format, and supports UTF-8, UTF-16 and UTF-32 character sets. The handler accepts an array of events (even if there is only one event, the event has to be sent in an array) and converts them to a Flume event based on the encoding specified in the request. If no encoding is specified, UTF-8 is assumed. The JSON handler supports UTF-8, UTF-16 and UTF-32. Events are represented as follows.

```
[{
  "headers" : {
    "timestamp" : "434324343",
    "host" : "random_host.example.com"
  },
  "body" : "random_body"
},
{
  "headers" : {
    "namenode" : "namenode.example.com",
    "datanode" : "random_datanode.example.com"
  },
  "body" : "really_random_body"
}]
```

To set the charset, the request must have content type specified as `application/json; charset=UTF-8` (replace UTF-8 with UTF-16 or UTF-32 as required).

One way to create an event in the format expected by this handler is to use `JSONEvent` provided in the Flume SDK and use Google Gson to create the JSON string using the `Gson#fromJson(Object, Type)` method. The type token to pass as the 2nd argument of this method for list of events can be created by:

```
Type type = new TypeToken<List<JSONEvent>>() {}.getType();
```

BlobHandler

By default `HTTPSource` splits JSON input into Flume events. As an alternative, `BlobHandler` is a handler for `HTTPSource` that returns an event that contains the request parameters as well as the Binary Large Object (BLOB) uploaded with this request. For example a PDF or JPG file. Note that this approach is not suitable for very large objects because it buffers up the entire BLOB in RAM.

| Property Name | Default | Description |
|-----------------------|---------|---|
| handler | – | The FQCN of this class: <code>org.apache.flume.sink.solr.morphline.BlobHandler</code> |
| handler.maxBlobLength | 1000000 | The maximum number of bytes to read and buffer for a given request |

Stress Source

StressSource is an internal load-generating source implementation which is very useful for stress tests. It allows User to configure the size of Event payload, with empty headers. User can configure total number of events to be sent as well maximum number of Successful Event to be delivered.

Required properties are in **bold**.

| Property Name | Default | Description |
|---------------------|---------|--|
| type | – | The component type name, needs to be <code>org.apache.flume.source.StressSource</code> |
| size | 500 | Payload size of each Event. Unit: byte |
| maxTotalEvents | -1 | Maximum number of Events to be sent |
| maxSuccessfulEvents | -1 | Maximum number of Events successfully sent |
| batchSize | 1 | Number of Events to be sent in one batch |

Example for agent named **a1**:

```
a1.sources = stresssource-1
a1.channels = memoryChannel-1
a1.sources.stresssource-1.type =
org.apache.flume.source.StressSource
a1.sources.stresssource-1.size = 10240
a1.sources.stresssource-1.maxTotalEvents = 1000000
a1.sources.stresssource-1.channels = memoryChannel-1
```

Legacy Sources

The legacy sources allow a Flume 1.x agent to receive events from Flume 0.9.4 agents. It accepts events in the Flume 0.9.4 format, converts them to the Flume 1.0 format, and stores them in the connected channel. The 0.9.4 event properties like timestamp, pri, host, nanos, etc get converted to 1.x event header attributes. The legacy source supports both Avro and Thrift RPC connections. To use this bridge between two Flume versions, you need to start a Flume 1.x agent with the avroLegacy or thriftLegacy source. The 0.9.4 agent should have the agent Sink pointing to the host/port of the 1.x agent.

Note

The reliability semantics of Flume 1.x are different from that of Flume 0.9.x. The E2E or DFO mode of a Flume 0.9.x agent will not be supported by the legacy source. The only supported 0.9.x mode is the best effort, though the reliability setting of the 1.x flow will be applicable to the events once they are saved into the Flume 1.x channel by the legacy source.

Required properties are in **bold**.

Avro Legacy Source

| Property Name | Default | Description |
|-----------------|-------------|---|
| channels | – | |
| type | – | The component type name, needs to be <code>org.apache.flume.source.avroLegacy.AvroLegacySource</code> |
| host | – | The hostname or IP address to bind to |
| port | – | The port # to listen on |
| selector.type | | replicating or multiplexing |
| selector.* | replicating | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type =
org.apache.flume.source.avroLegacy.AvroLegacySource
a1.sources.r1.host = 0.0.0.0
a1.sources.r1.bind = 6666
a1.sources.r1.channels = c1
```

Thrift Legacy Source

| Property Name | Default | Description |
|----------------|-------------|---|
| channels | – | |
| type | – | The component type name, needs to be <code>org.apache.flume.source.thriftLegacy.ThriftLegacySource</code> |
| host | – | The hostname or IP address to bind to |
| port | – | The port # to listen on |
| selector.type | | replicating or multiplexing |
| selector.* | replicating | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type =
org.apache.flume.source.thriftLegacy.ThriftLegacySource
a1.sources.r1.host = 0.0.0.0
a1.sources.r1.bind = 6666
a1.sources.r1.channels = c1
```

Custom Source

A custom source is your own implementation of the Source interface. A custom source's class and its dependencies must be included in the agent's classpath when starting the Flume agent. The type of the custom source is its FQCN.

| Property Name | Default | Description |
|----------------|-------------|--|
| channels | – | |
| type | – | The component type name, needs to be your FQCN |
| selector.type | | replicating or multiplexing |
| selector.* | replicating | Depends on the selector.type value |
| interceptors | – | Space-separated list of interceptors |
| interceptors.* | | |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = org.example.MySource
a1.sources.r1.channels = c1
```

Scribe Source

Scribe is another type of ingest system. To adopt existing Scribe ingest system, Flume should use ScribeSource based on Thrift with compatible transferring protocol. For deployment of Scribe please follow the guide from Facebook. Required properties are in **bold**.

| Property Name | Default | Description |
|---------------------|----------|--|
| type | – | The component type name, needs to be org.apache.flume.source.scribe.ScribeSource |
| port | 1499 | Port that Scribe should be connected |
| maxReadBuffer Bytes | 16384000 | Thrift Default FrameBuffer Size |
| workerThreads | 5 | Handling threads number in Thrift |
| selector.type | | |
| selector.* | | |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.type = org.apache.flume.source.scribe.ScribeSource
a1.sources.r1.port = 1463
```



```
a1.sources.r1.workerThreads = 5  
a1.sources.r1.channels = c1
```

Flume Sinks

HDFS Sink

This sink writes events into the Hadoop Distributed File System (HDFS). It currently supports creating text and sequence files. It supports compression in both file types. The files can be rolled (close current file and create a new one) periodically based on the elapsed time or size of data or number of events. It also buckets/partitions data by attributes like timestamp or machine where the event originated. The HDFS directory path may contain formatting escape sequences that will be replaced by the HDFS sink to generate a directory/file name to store the events. Using this sink requires Hadoop to be installed so that Flume can use the Hadoop jars to communicate with the HDFS cluster. Note that a version of Hadoop that supports the `sync()` call is required.

The following are the escape sequences supported:

| Alias | Description |
|---------|--|
| %{host} | Substitute value of event header named "host". Arbitrary header names are supported. |
| %t | Unix time in milliseconds |
| %a | locale's short weekday name (Mon, Tue, ...) |
| %A | locale's full weekday name (Monday, Tuesday, ...) |
| %b | locale's short month name (Jan, Feb, ...) |
| %B | locale's long month name (January, February, ...) |
| %c | locale's date and time (Thu Mar 3 23:05:25 2005) |
| %d | day of month (01) |
| %e | day of month without padding (1) |
| %D | date; same as %m/%d/%y |
| %H | hour (00..23) |
| %I | hour (01..12) |
| %j | day of year (001..366) |
| %k | hour (0..23) |
| %m | month (01..12) |
| %n | month without padding (1..12) |
| %M | minute (00..59) |

| Alias | Description |
|-------|---|
| %p | locale's equivalent of am or pm |
| %s | seconds since 1970-01-01 00:00:00 UTC |
| %S | second (00..59) |
| %y | last two digits of year (00..99) |
| %Y | year (2010) |
| %z | +hhmm numeric timezone (for example, -0400) |

The file in use will have the name mangled to include ".tmp" at the end. Once the file is closed, this extension is removed. This allows excluding partially complete files in the directory. Required properties are in **bold**.

Note

For all of the time related escape sequences, a header with the key "timestamp" must exist among the headers of the event (unless `hdfs.useLocalTimeStamp` is set to `true`). One way to add this automatically is to use the `TimestampInterceptor`.

| Name | Default | Description |
|--------------------------------|-------------------|--|
| channel | – | |
| type | – | The component type name, needs to be <code>hdfs</code> |
| hdfs.path | – | HDFS directory path (eg <code>hdfs://namenode/flume/webdata/</code>) |
| <code>hdfs.filePrefix</code> | FlumeData | Name prefixed to files created by Flume in <code>hdfs</code> directory |
| <code>hdfs.fileSuffix</code> | – | Suffix to append to file (eg <code>.avro</code> - <i>NOTE: period is not automatically added</i>) |
| <code>hdfs.inUsePrefix</code> | – | Prefix that is used for temporal files that flume actively writes into |
| <code>hdfs.inUseSuffix</code> | <code>.tmp</code> | Suffix that is used for temporal files that flume actively writes into |
| <code>hdfs.rollInterval</code> | 30 | Number of seconds to wait before rolling current file (0 = never roll based on time interval) |
| <code>hdfs.rollSize</code> | 1024 | File size to trigger roll, in bytes (0: never roll based on file size) |
| <code>hdfs.rollCount</code> | 10 | Number of events written to file before it rolled (0 = never roll based on number of events) |
| <code>hdfs.idleTimeout</code> | 0 | Timeout after which inactive files get closed (0 = disable automatic closing of idle files) |

| Name | Default | Description |
|------------------------|--------------|--|
| hdfs.batchSize | 100 | number of events written to file before it is flushed to HDFS |
| hdfs.codeC | – | Compression codec. one of following : gzip, bzip2, lzo, lzop, snappy |
| hdfs.fileType | SequenceFile | File format: currently SequenceFile, DataStream or CompressedStream (1)DataStream will not compress output file and please don't set codeC (2)CompressedStream requires set hdfs.codeC with an available codeC |
| hdfs.maxOpenFiles | 5000 | Allow only this number of open files. If this number is exceeded, the oldest file is closed. |
| hdfs.minBlockReplicas | – | Specify minimum number of replicas per HDFS block. If not specified, it comes from the default Hadoop config in the classpath. |
| hdfs.writeFormat | – | Format for sequence file records. One of “Text” or “Writable” (the default). |
| hdfs.callTimeout | 10000 | Number of milliseconds allowed for HDFS operations, such as open, write, flush, close. This number should be increased if many HDFS timeout operations are occurring. |
| hdfs.threadsPoolSize | 10 | Number of threads per HDFS sink for HDFS IO ops (open, write, etc.) |
| hdfs.rollTimerPoolSize | 1 | Number of threads per HDFS sink for scheduling timed file rolling |
| hdfs.kerberosPrincipal | – | Kerberos user principal for accessing secure HDFS |
| hdfs.kerberosKeytab | – | Kerberos keytab for accessing secure HDFS |
| hdfs.proxyUser | | |
| hdfs.round | false | Should the timestamp be rounded down (if true, affects all time based escape sequences except %t) |
| hdfs.roundValue | 1 | Rounded down to the highest multiple of this (in the unit configured using <code>hdfs.roundUnit</code>), less than current time. |
| hdfs.roundUnit | second | The unit of the round down value - second, minute or hour. |
| hdfs.timeZone | Local Time | Name of the timezone that should be used for resolving the directory path, e.g. America/Los_Angeles. |

| Name | Default | Description |
|------------------------|---------|--|
| hdfs.useLocalTimeStamp | false | Use the local time (instead of the timestamp from the event header) while replacing the escape sequences. |
| hdfs.closeTries | 0 | Number of times the sink must try renaming a file, after initiating a close attempt. If set to 1, this sink will not re-try a failed rename (due to, for example, NameNode or DataNode failure), and may leave the file in an open state with a .tmp extension. If set to 0, the sink will try to rename the file until the file is eventually renamed (there is no limit on the number of times it would try). The file may still remain open if the close call fails but the data will be intact and in this case, the file will be closed only after a Flume restart. |
| hdfs.retryInterval | 180 | Time in seconds between consecutive attempts to close a file. Each close call costs multiple RPC round-trips to the Namenode, so setting this too low can cause a lot of load on the name node. If set to 0 or less, the sink will not attempt to close the file if the first attempt fails, and may leave the file open or with a ".tmp" extension. |
| serializer | TEXT | Other possible options include <code>avro_event</code> or the fully-qualified class name of an implementation of the <code>EventSerializer.Builder</code> interface. |
| serializer.* | | |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = hdfs
a1.sinks.k1.channel = c1
a1.sinks.k1.hdfs.path = /flume/events/%y-%m-%d/%H%M/%S
a1.sinks.k1.hdfs.filePrefix = events-
a1.sinks.k1.hdfs.round = true
a1.sinks.k1.hdfs.roundValue = 10
a1.sinks.k1.hdfs.roundUnit = minute
```

The above configuration will round down the timestamp to the last 10th minute. For example, an event with timestamp 11:54:34 AM, June 12, 2012 will cause the hdfs path to become `/flume/events/2012-06-12/1150/00`.

Hive Sink

This sink streams events containing delimited text or JSON data directly into a Hive table or partition. Events are written using Hive transactions. As soon as a set of events are committed to Hive, they become immediately visible to Hive queries. Partitions to which flume will stream to can either be pre-created or, optionally, Flume can create them if they are missing. Fields from incoming event data are mapped to corresponding columns in the Hive table. **This sink is provided as a preview feature and not recommended for use in production.**

| Name | Default | Description |
|----------------------|---------|---|
| channel | – | |
| type | – | The component type name, needs to be <code>hive</code> |
| hive.metastore | – | Hive metastore URI (eg <code>thrift://a.b.com:9083</code>) |
| hive.database | – | Hive database name |
| hive.table | – | Hive table name |
| hive.partition | – | Comma separate list of partition values identifying the partition to write to. May contain escape sequences. E.g: If the table is partitioned by (continent: string, country :string, time : string) then ‘Asia,India,2014-02-26-01-21’ will indicate continent=Asia,country=India,time=2014-02-26-01-21 |
| hive.txnsPerBatchAsk | 100 | Hive grants a <i>batch of transactions</i> instead of single transactions to streaming clients like Flume. This setting configures the number of desired transactions per Transaction Batch. Data from all transactions in a single batch end up in a single file. Flume will write a maximum of batchSize events in each transaction in the batch. This setting in conjunction with batchSize provides control over the size of each file. Note that eventually Hive will transparently compact these files into larger files. |
| heartBeatInterval | 240 | (In seconds) Interval between consecutive heartbeats sent to Hive to keep unused transactions from expiring. Set this value to 0 to disable heartbeats. |

| Name | Default | Description |
|----------------------|------------|--|
| autoCreatePartitions | true | Flume will automatically create the necessary Hive partitions to stream to |
| batchSize | 15000 | Max number of events written to Hive in a single Hive transaction |
| maxOpenConnections | 500 | Allow only this number of open connections. If this number is exceeded, the least recently used connection is closed. |
| callTimeout | 10000 | (In milliseconds) Timeout for Hive & HDFS I/O operations, such as openTxn, write, commit, abort. |
| serializer | | Serializer is responsible for parsing out field from the event and mapping them to columns in the hive table. Choice of serializer depends upon the format of the data in the event. Supported serializers: DELIMITED and JSON |
| roundUnit | minute | The unit of the round down value - second, minute or hour. |
| roundValue | 1 | Rounded down to the highest multiple of this (in the unit configured using hive.roundUnit), less than current time |
| timeZone | Local Time | Name of the timezone that should be used for resolving the escape sequences in partition, e.g. America/Los_Angeles. |
| useLocalTimeStamp | false | Use the local time (instead of the timestamp from the event header) while replacing the escape sequences. |

Following serializers are provided for Hive sink:

JSON: Handles UTF8 encoded Json (strict syntax) events and requires no configuration. Object names in the JSON are mapped directly to columns with the same name in the Hive table. Internally uses org.apache.hive.hcatalog.data.JsonSerDe but is independent of the Serde of the Hive table. This serializer requires HCatalog to be installed.

DELIMITED: Handles simple delimited textual events. Internally uses LazySimpleSerde but is independent of the Serde of the Hive table.

| Name | Default | Description |
|------------------------------|---------|---|
| serializer.delimiter | , | (Type: string) The field delimiter in the incoming data. To use special characters, surround them with double quotes like “\t” |
| serializer.fieldnames | – | The mapping from input fields to columns in hive table. Specified as a comma separated list (no spaces) of hive table columns names, identifying the input fields in order of their occurrence. To skip fields leave the column name unspecified. Eg. ‘time,,ip,message’ indicates the 1st, 3rd and 4th fields in input map to time, ip and message columns in the hive table. |
| serializer.serdeSeparator | Ctrl-A | (Type: character) Customizes the separator used by underlying serde. There can be a gain in efficiency if the fields in serializer.fieldnames are in same order as table columns, the serializer.delimiter is same as the serializer.serdeSeparator and number of fields in serializer.fieldnames is less than or equal to number of table columns, as the fields in incoming event body do not need to be reordered to match order of table columns. Use single quotes for special characters like ‘\t’. Ensure input fields do not contain this character. NOTE: If serializer.delimiter is a single character, preferably set this to the same character |

The following are the escape sequences supported:

| Alias | Description |
|---------|--|
| %{host} | Substitute value of event header named “host”. Arbitrary header names are supported. |
| %t | Unix time in milliseconds |
| %a | locale’s short weekday name (Mon, Tue, ...) |
| %A | locale’s full weekday name (Monday, Tuesday, ...) |
| %b | locale’s short month name (Jan, Feb, ...) |
| %B | locale’s long month name (January, February, ...) |
| %c | locale’s date and time (Thu Mar 3 23:05:25 2005) |
| %d | day of month (01) |
| %D | date; same as %m/%d/%y |

| Alias | Description |
|-------|---|
| %H | hour (00..23) |
| %I | hour (01..12) |
| %j | day of year (001..366) |
| %k | hour (0..23) |
| %m | month (01..12) |
| %M | minute (00..59) |
| %p | locale's equivalent of am or pm |
| %s | seconds since 1970-01-01 00:00:00 UTC |
| %S | second (00..59) |
| %y | last two digits of year (00..99) |
| %Y | year (2010) |
| %z | +hhmm numeric timezone (for example, -0400) |

Note

For all of the time related escape sequences, a header with the key “timestamp” must exist among the headers of the event (unless `useLocalTimeStamp` is set to `true`). One way to add this automatically is to use the `TimestampInterceptor`.

Example Hive table :

```
create table weblogs ( id int , msg string )
    partitioned by (continent string, country string, time string)
    clustered by (id) into 5 buckets
    stored as orc;
```

Example for agent named a1:

```
a1.channels = c1
a1.channels.c1.type = memory
a1.sinks = k1
a1.sinks.k1.type = hive
a1.sinks.k1.channel = c1
a1.sinks.k1.hive.metastore = thrift://127.0.0.1:9083
a1.sinks.k1.hive.database = logsdb
a1.sinks.k1.hive.table = weblogs
a1.sinks.k1.hive.partition = asia,{country},%y-%m-%d-%H-%M
a1.sinks.k1.useLocalTimeStamp = false
a1.sinks.k1.round = true
a1.sinks.k1.roundValue = 10
a1.sinks.k1.roundUnit = minute
a1.sinks.k1.serializer = DELIMITED
```



```
a1.sinks.k1.serializer.delimiter = "\\t"  
a1.sinks.k1.serializer.serdeSeparator = '\\t'  
a1.sinks.k1.serializer.fieldnames = id,msg
```

The above configuration will round down the timestamp to the last 10th minute. For example, an event with timestamp header set to 11:54:34 AM, June 12, 2012 and 'country' header set to 'india' will evaluate to the partition (continent='asia',country='india',time='2012-06-12-11-50'. The serializer is configured to accept tab separated input containing three fields and to skip the second field.

Logger Sink

Logs event at INFO level. Typically useful for testing/debugging purpose. Required properties are in **bold**.

| Property Name | Default | Description |
|----------------------|---------|--|
| channel | — | |
| type | — | The component type name, needs to be <code>logger</code> |
| maxBytesToLog | 16 | Maximum number of bytes of the Event body to log |

Example for agent named a1:

```
a1.channels = c1  
a1.sinks = k1  
a1.sinks.k1.type = logger  
a1.sinks.k1.channel = c1
```

Avro Sink

This sink forms one half of Flume's tiered collection support. Flume events sent to this sink are turned into Avro events and sent to the configured hostname / port pair. The events are taken from the configured Channel in batches of the configured batch size. Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------|---------|--|
| channel | — | |
| type | — | The component type name, needs to be <code>avro</code> . |
| hostname | — | The hostname or IP address to bind to. |

| Property Name | Default | Description |
|---------------------------|---------|---|
| port | – | The port # to listen on. |
| batch-size | 100 | number of event to batch together for send. |
| connect-timeout | 20000 | Amount of time (ms) to allow for the first (handshake) request. |
| request-timeout | 20000 | Amount of time (ms) to allow for requests after the first. |
| reset-connection-interval | none | Amount of time (s) before the connection to the next hop is reset. This will force the Avro Sink to reconnect to the next hop. This will allow the sink to connect to hosts behind a hardware load-balancer when new hosts are added without having to restart the agent. |
| compression-type | none | This can be “none” or “deflate”. The compression-type must match the compression-type of matching AvroSource |
| compression-level | 6 | The level of compression to compress event. 0 = no compression and 1-9 is compression. The higher the number the more compression |
| ssl | false | Set to true to enable SSL for this AvroSink. When configuring SSL, you can optionally set a “truststore”, “truststore-password”, “truststore-type”, and specify whether to “trust-all-certs”. |
| trust-all-certs | false | If this is set to true, SSL server certificates for remote servers (Avro Sources) will not be checked. This should NOT be used in production because it makes it easier for an attacker to execute a man-in-the-middle attack and “listen in” on the encrypted connection. |
| truststore | – | The path to a custom Java truststore file. Flume uses the certificate authority information in this file to determine whether the remote Avro Source’s SSL authentication credentials should be trusted. If not specified, the default Java JSSE certificate authority files (typically “jssecacerts” or “cacerts” in the Oracle JRE) will be used. |
| truststore-password | – | The password for the specified truststore. |
| truststore-type | JKS | The type of the Java truststore. This can be “JKS” or other supported Java truststore type. |

| Property Name | Default | Description |
|-------------------|---|---|
| exclude-protocols | SSLv3 | Space-separated list of SSL/TLS protocols to exclude. SSLv3 will always be excluded in addition to the protocols specified. |
| maxIoWorkers | 2 * the number of available processors in the machine | The maximum number of I/O worker threads. This is configured on the NettyAvroRpcClient NioClientSocketChannelFactory. |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = avro
a1.sinks.k1.channel = c1
a1.sinks.k1.hostname = 10.10.10.10
a1.sinks.k1.port = 4545
```

Thrift Sink

This sink forms one half of Flume's tiered collection support. Flume events sent to this sink are turned into Thrift events and sent to the configured hostname / port pair. The events are taken from the configured Channel in batches of the configured batch size.

Thrift sink can be configured to start in secure mode by enabling kerberos authentication. To communicate with a Thrift source started in secure mode, the Thrift sink should also operate in secure mode. **client-principal** and **client-keytab** are the properties used by the Thrift sink to authenticate to the kerberos KDC. The **server-principal** represents the principal of the Thrift source this sink is configured to connect to in secure mode. Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------|---------|--|
| channel | — | |
| type | — | The component type name, needs to be <code>thrift</code> . |
| hostname | — | The hostname or IP address to bind to. |
| port | — | The port # to listen on. |

| Property Name | Default | Description |
|---------------------------|---------|---|
| batch-size | 100 | number of event to batch together for send. |
| connect-timeout | 20000 | Amount of time (ms) to allow for the first (handshake) request. |
| request-timeout | 20000 | Amount of time (ms) to allow for requests after the first. |
| connection-reset-interval | none | Amount of time (s) before the connection to the next hop is reset. This will force the Thrift Sink to reconnect to the next hop. This will allow the sink to connect to hosts behind a hardware load-balancer when new hosts are added without having to restart the agent. |
| ssl | false | Set to true to enable SSL for this ThriftSink. When configuring SSL, you can optionally set a “truststore”, “truststore-password” and “truststore-type” |
| truststore | — | The path to a custom Java truststore file. Flume uses the certificate authority information in this file to determine whether the remote Thrift Source’s SSL authentication credentials should be trusted. If not specified, the default Java JSSE certificate authority files (typically “jssecacerts” or “cacerts” in the Oracle JRE) will be used. |
| truststore-password | — | The password for the specified truststore. |
| truststore-type | JKS | The type of the Java truststore. This can be “JKS” or other supported Java truststore type. |
| exclude-protocols | SSLv3 | Space-separated list of SSL/TLS protocols to exclude |
| kerberos | false | Set to true to enable kerberos authentication. In kerberos mode, client-principal, client-keytab and server-principal are required for successful authentication and communication to a kerberos enabled Thrift Source. |
| client-principal | — | The kerberos principal used by the Thrift Sink to authenticate to the kerberos KDC. |
| client-keytab | — | The keytab location used by the Thrift Sink in combination with the client-principal to authenticate to the kerberos KDC. |
| server-principal | — | The kerberos principal of the Thrift Source to which the Thrift Sink is configured to connect to. |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = thrift
a1.sinks.k1.channel = c1
a1.sinks.k1.hostname = 10.10.10.10
a1.sinks.k1.port = 4545
```

IRC Sink

The IRC sink takes messages from attached channel and relays those to configured IRC destinations. Required properties are in **bold**.

| Property Name | Default | Description |
|-------------------|---------|--|
| channel | – | |
| type | – | The component type name, needs to be <code>irc</code> |
| hostname | – | The hostname or IP address to connect to |
| port | 6667 | The port number of remote host to connect |
| nick | – | Nick name |
| user | – | User name |
| password | – | User password |
| chan | – | channel |
| name | | |
| splitlines | – | (boolean) |
| splitchars | n | line separator (if you were to enter the default value into the config file, then you would need to escape the backslash, like this: “\n”) |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = irc
a1.sinks.k1.channel = c1
a1.sinks.k1.hostname = irc.yourdomain.com
a1.sinks.k1.nick = flume
a1.sinks.k1.chan = #flume
```

File Roll Sink

Stores events on the local filesystem. Required properties are in **bold**.

| Property Name | Default | Description |
|--------------------------------|---------|--|
| channel | – | |
| type | – | The component type name, needs to be <code>file_roll</code> . |
| sink.directory | – | The directory where files will be stored |
| <code>sink.rollInterval</code> | 30 | Roll the file every 30 seconds. Specifying 0 will disable rolling and cause all events to be written to a single file. |
| <code>sink.serializer</code> | TEXT | Other possible options include <code>avro_event</code> or the FQCN of an implementation of <code>EventSerializer.Builder</code> interface. |
| <code>batchSize</code> | 100 | |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = file_roll
a1.sinks.k1.channel = c1
a1.sinks.k1.sink.directory = /var/log/flume
```

Null Sink

Discards all events it receives from the channel. Required properties are in **bold**.

| Property Name | Default | Description |
|------------------------|---------|--|
| channel | – | |
| type | – | The component type name, needs to be <code>null</code> . |
| <code>batchSize</code> | 100 | |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = null
a1.sinks.k1.channel = c1
```

HBaseSinks

HBaseSink

This sink writes data to HBase. The Hbase configuration is picked up from the first hbase-site.xml encountered in the classpath. A class implementing HbaseEventSerializer which is specified by the configuration is used to convert the events into HBase puts and/or increments. These puts and increments are then written to HBase. This sink provides the same consistency guarantees as HBase, which is currently row-wise atomicity. In the event of Hbase failing to write certain events, the sink will replay all events in that transaction.

The HBaseSink supports writing data to secure HBase. To write to secure HBase, the user the agent is running as must have write permissions to the table the sink is configured to write to. The principal and keytab to use to authenticate against the KDC can be specified in the configuration. The hbase-site.xml in the Flume agent's classpath must have authentication set to `kerberos` (For details on how to do this, please refer to HBase documentation).

For convenience, two serializers are provided with Flume. The SimpleHbaseEventSerializer (org.apache.flume.sink.hbase.SimpleHbaseEventSerializer) writes the event body as-is to HBase, and optionally increments a column in Hbase. This is primarily an example implementation. The RegexHbaseEventSerializer (org.apache.flume.sink.hbase.RegexHbaseEventSerializer) breaks the event body based on the given regex and writes each part into different columns.

The type is the FQCN: org.apache.flume.sink.hbase.HBaseSink.

Required properties are in **bold**.

| Property Name | Default | Description |
|------------------------|---------|---|
| channel | — | |
| type | — | The component type name, needs to be <code>hbase</code> |
| table | — | The name of the table in Hbase to write to. |
| columnFamily | — | The column family in Hbase to write to. |
| zookeeperQuorum | — | The quorum spec. This is the value for the property <code>hbase.zookee</code> |

| Property Name | Default | Description |
|--------------------|--|---|
| | | per.quorum in hbase-site.xml |
| znodeParent | /hbase | The base path for the znode for the -ROOT- region. Value of zookeeper.znode.parent in hbase-site.xml |
| batchSize | 100 | Number of events to be written per txn. |
| coalesceIncrements | false | Should the sink coalesce multiple increments to a cell per batch. This might give better performance if there are multiple increments to a limited number of cells. |
| serializer | org.apache.flume.sink.hbase.SimpleHbaseEventSerializer | Default increment column = "iCol", payload column = "pCol". |
| serializer.* | – | Properties to be passed to the serializer. |
| kerberosPrincipal | – | Kerberos user principal for accessing secure HBase |
| kerberosKeytab | – | Kerberos keytab for accessing secure HBase |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = hbase
a1.sinks.k1.table = foo_table
a1.sinks.k1.columnFamily = bar_cf
a1.sinks.k1.serializer =
org.apache.flume.sink.hbase.RegexHbaseEventSerializer
a1.sinks.k1.channel = c1
```

AsyncHBaseSink

This sink writes data to HBase using an asynchronous model. A class implementing AsyncHbaseEventSerializer which is specified by the configuration is used to

convert the events into HBase puts and/or increments. These puts and increments are then written to HBase. This sink uses the [Asynchbase API](#) to write to HBase. This sink provides the same consistency guarantees as HBase, which is currently row-wise atomicity. In the event of Hbase failing to write certain events, the sink will replay all events in that transaction. The type is the FQCN: org.apache.flume.sink.hbase.AsyncHBaseSink. Required properties are in **bold**.

| Property Name | Default | Description |
|---------------------|---------|---|
| channel | – | |
| type | – | The component type name, needs to be <code>asynchbase</code> |
| table | – | The name of the table in Hbase to write to. |
| zookeeper Quorum | – | The quorum spec. This is the value for the property <code>hbase.zookeeper.quorum</code> in <code>hbase-site.xml</code> |
| znodeParent | /hbase | The base path for the znode for the -ROOT- region. Value of <code>zookeeper.znode.parent</code> in <code>hbase-site.xml</code> |
| columnFamily | – | The column family in Hbase to write to. |
| batchSize | 100 | Number of events to be written per txn. |
| coalesceIncrements | false | Should the sink coalesce multiple increments to a cell per batch. This might give better performance if there are multiple increments to a limited number of cells. |
| timeout | 60000 | The length of time (in milliseconds) the sink waits for acks from hbase for all events in a transaction. |

| Property Name | Default | Description |
|---------------|---|--|
| serializer | org.apache.flume.sink.hbase.SimpleAsyncHbaseEventSerializer | |
| serializer.* | – | Properties to be passed to the serializer. |

Note that this sink takes the Zookeeper Quorum and parent znode information in the configuration. Zookeeper Quorum and parent node configuration may be specified in the flume configuration file. Alternatively, these configuration values are taken from the first hbase-site.xml file in the classpath.

If these are not provided in the configuration, then the sink will read this information from the first hbase-site.xml file in the classpath.

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = asynchbase
a1.sinks.k1.table = foo_table
a1.sinks.k1.columnFamily = bar_cf
a1.sinks.k1.serializer =
org.apache.flume.sink.hbase.SimpleAsyncHbaseEventSerializer
a1.sinks.k1.channel = c1
```

MorphlineSolrSink

This sink extracts data from Flume events, transforms it, and loads it in near-real-time into Apache Solr servers, which in turn serve queries to end users or search applications.

This sink is well suited for use cases that stream raw data into HDFS (via the HdfsSink) and simultaneously extract, transform and load the same data into Solr (via MorphlineSolrSink). In particular, this sink can process arbitrary heterogeneous raw data from disparate data sources and turn it into a data model that is useful to Search applications.

The ETL functionality is customizable using a [morphline configuration file](#) that defines a chain of transformation commands that pipe event records from one command to another.

Morphlines can be seen as an evolution of Unix pipelines where the data model is generalized to work with streams of generic records, including arbitrary binary payloads. A morphline command is a bit like a Flume Interceptor. Morphlines can be embedded into Hadoop components such as Flume.

Commands to parse and transform a set of standard data formats such as log files, Avro, CSV, Text, HTML, XML, PDF, Word, Excel, etc. are provided out of the box, and additional custom commands and parsers for additional data formats can be added as morphline plugins. Any kind of data format can be indexed and any Solr documents for any kind of Solr schema can be generated, and any custom ETL logic can be registered and executed.

Morphlines manipulate continuous streams of records. The data model can be described as follows: A record is a set of named fields where each field has an ordered list of one or more values. A value can be any Java Object. That is, a record is essentially a hash table where each hash table entry contains a String key and a list of Java Objects as values. (The implementation uses Guava's `ArrayListMultimap`, which is a `ListMultimap`). Note that a field can have multiple values and any two records need not use common field names.

This sink fills the body of the Flume event into the `attachment_body` field of the morphline record, as well as copies the headers of the Flume event into record fields of the same name. The commands can then act on this data.

Routing to a SolrCloud cluster is supported to improve scalability. Indexing load can be spread across a large number of `MorphlineSolrSinks` for improved scalability. Indexing load can be replicated across multiple `MorphlineSolrSinks` for high availability, for example using Flume features such as Load balancing Sink Processor. `MorphlineInterceptor` can also help to implement dynamic routing to multiple Solr collections (e.g. for multi-tenancy).

The morphline and solr jars required for your environment must be placed in the lib directory of the Apache Flume installation.

The type is the FQCN: `org.apache.flume.sink.solr.morphline.MorphlineSolrSink`

Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------------------|--|--|
| channel | – | |
| type | – | The component type name, needs to be <code>org.apache.flume.sink.solr.morphline.MorphlineSolrSink</code> |
| morphlineFile | – | The relative or absolute path on the local file system to the morphline configuration file. Example: <code>/etc/flume-ng/conf/morphline.conf</code> |
| morphlineId | null | Optional name used to identify a morphline if there are multiple morphlines in a morphline config file |
| batchSize | 1000 | The maximum number of events to take per flume transaction. |
| batchDurationMillis | 1000 | The maximum duration per flume transaction (ms). The transaction commits after this duration or when batchSize is exceeded, whichever comes first. |
| handlerClass | <code>org.apache.flume.sink.solr.morphline.MorphlineHandlerImpl</code> | The FQCN of a class implementing <code>org.apache.flume.sink.solr.morphline.MorphlineHandler</code> |
| isProductionMode | false | This flag should be enabled for mission critical, large-scale online production systems that need to make progress without downtime when unrecoverable exceptions occur. Corrupt or malformed parser input data, parser bugs, and errors related to unknown Solr schema fields produce unrecoverable exceptions. |
| recoverableExceptionClasses | <code>org.apache.solr.client.solrj.SolrServerException</code> | Comma separated list of recoverable exceptions that tend to be transient, in which case the corresponding task can be retried. Examples include network |

| Property Name | Default | Description |
|---------------------------------|---------|--|
| | | connection errors, timeouts, etc. When the production mode flag is set to true, the recoverable exceptions configured using this parameter will not be ignored and hence will lead to retries. |
| isIgnoringRecoverableExceptions | false | This flag should be enabled, if an unrecoverable exception is accidentally misclassified as recoverable. This enables the sink to make progress and avoid retrying an event forever. |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type =
org.apache.flume.sink.solr.morphline.MorphlineSolrSink
a1.sinks.k1.channel = c1
a1.sinks.k1.morphlineFile = /etc/flume-ng/conf/morphline.conf
# a1.sinks.k1.morphlineId = morphline1
# a1.sinks.k1.batchSize = 1000
# a1.sinks.k1.batchDurationMillis = 1000
```

ElasticSearchSink

This sink writes data to an elasticsearch cluster. By default, events will be written so that the [Kibana](#) graphical interface can display them - just as if [logstash](#) wrote them.

The elasticsearch and lucene-core jars required for your environment must be placed in the lib directory of the Apache Flume installation. Elasticsearch requires that the major version of the client JAR match that of the server and that both are running the same minor version of the JVM. SerializationExceptions will appear if this is incorrect. To select the required version first determine the version of elasticsearch and the JVM version the target cluster is running. Then select an elasticsearch client library which matches the major version. A 0.19.x client can talk to a 0.19.x cluster; 0.20.x can talk to 0.20.x and 0.90.x can talk to 0.90.x. Once the elasticsearch version has been determined then read the pom.xml file to determine the correct lucene-core

JAR version to use. The Flume agent which is running the ElasticSearchSink should also match the JVM the target cluster is running down to the minor version.

Events will be written to a new index every day. The name will be <indexName>-yyyy-MM-dd where <indexName> is the indexName parameter. The sink will start writing to a new index at midnight UTC.

Events are serialized for elasticsearch by the ElasticSearchLogStashEventSerializer by default. This behaviour can be overridden with the serializer parameter. This parameter accepts implementations of org.apache.flume.sink.elasticsearch.ElasticSearchEventSerializer or org.apache.flume.sink.elasticsearch.ElasticSearchIndexRequestBuilderFactory. Implementing ElasticSearchEventSerializer is deprecated in favour of the more powerful ElasticSearchIndexRequestBuilderFactory.

The type is the FQCN: org.apache.flume.sink.elasticsearch.ElasticSearchSink

Required properties are in **bold**.

| Property Name | Default | Description |
|-------------------|---------|---|
| channel | – | |
| type | – | The component type name, needs to be org.apache.flume.sink.elasticsearch.ElasticSearchSink |
| host Names | – | Comma separated list of hostname:port, if the port is not present the default port '9300' will be used |
| index Name | flume | The name of the index which the date will be appended to. Example 'flume' -> 'flume-yyyy-MM-dd' Arbitrary header substitution is supported, eg. %{header} replaces with value of named event header |
| index Type | logs | The type to index the document to, defaults to 'log' Arbitrary header substitution is supported, eg. |

| Property Name | Default | Description |
|---------------|--|--|
| | | % {header} replaces with value of named event header |
| clusterName | elasticsearch | Name of the ElasticSearch cluster to connect to |
| batchSize | 100 | Number of events to be written per txn. |
| ttl | — | TTL in days, when set will cause the expired documents to be deleted automatically, if not set documents will never be automatically deleted. TTL is accepted both in the earlier form of integer only e.g. a1.sinks.k1.ttl = 5 and also with a qualifier ms (millisecond), s (second), m (minute), h (hour), d (day) and w (week). Example a1.sinks.k1.ttl = 5d will set TTL to 5 days. Follow http://www.elasticsearch.org/guide/reference/mapping/ttl-field/ for more information. |
| serializer | org.apache.flume.sink.elasticsearch.ElasticSearchLogStashEventSerializer | The ElasticSearchIndexRequestBuilderFactory or ElasticSearchEventSerializer to use. Implementations of either class are accepted but ElasticSearchIndexRequestBuilderFactory is preferred. |
| serializer.* | — | Properties to be passed to the serializer. |

Note

Header substitution is a handy to use the value of an event header to dynamically decide the indexName and indexType to use when storing the event. Caution should be used in using this feature as the event submitter now has control of the indexName

and `indexType`. Furthermore, if the elasticsearch REST client is used then the event submitter has control of the URL path used.

Example for agent named `a1`:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = elasticsearch
a1.sinks.k1.hostNames = 127.0.0.1:9200,127.0.0.2:9300
a1.sinks.k1.indexName = foo_index
a1.sinks.k1.indexType = bar_type
a1.sinks.k1.clusterName = foobar_cluster
a1.sinks.k1.batchSize = 500
a1.sinks.k1.ttl = 5d
a1.sinks.k1.serializer =
org.apache.flume.sink.elasticsearch.ElasticSearchDynamicSerializer
a1.sinks.k1.channel = c1
```

Kite Dataset Sink

Experimental sink that writes events to a [Kite Dataset](#). This sink will deserialize the body of each incoming event and store the resulting record in a Kite Dataset. It determines target Dataset by loading a dataset by URI.

The only supported serialization is avro, and the record schema must be passed in the event headers, using either `flume.avro.schema.literal` with the JSON schema representation or `flume.avro.schema.url` with a URL where the schema may be found (`hdfs:/...` URIs are supported). This is compatible with the Log4jAppender flume client and the spooling directory source's Avro deserializer using `deserializer.schemaType = LITERAL`.

Note 1: The `flume.avro.schema.hash` header is **not supported**. Note 2: In some cases, file rolling may occur slightly after the roll interval has been exceeded. However, this delay will not exceed 5 seconds. In most cases, the delay is negligible.

| Property Name | Default | Description |
|-------------------------|---------|---|
| channel | — | |
| type | — | Must be <code>org.apache.flume.sink.kite.DatasetSink</code> |
| kite.dataset.uri | — | URI of the dataset to open |

| Property Name | Default | Description |
|------------------------------|---------|--|
| kite.repo.uri | – | URI of the repository to open (deprecated; use kite.dataset.uri instead) |
| kite.dataset.namespace | – | Namespace of the Dataset where records will be written (deprecated; use kite.dataset.uri instead) |
| kite.dataset.name | – | Name of the Dataset where records will be written (deprecated; use kite.dataset.uri instead) |
| kite.batchSize | 100 | Number of records to process in each batch |
| kite.rollInterval | 30 | Maximum wait time (seconds) before data files are released |
| kite.flushable.commitOnBatch | true | If <code>true</code> , the Flume transaction will be committed and the writer will be flushed on each batch of <code>kite.batchSize</code> records. This setting only applies to flushable datasets. When <code>true</code> , it's possible for temp files with committed data to be left in the dataset directory. These files need to be recovered by hand for the data to be visible to DatasetReaders. |
| kite.syncable.syncOnBatch | true | Controls whether the sink will also sync data when committing the transaction. This setting only applies to syncable datasets. Syncing gaurentees that data will be written on stable storage on the remote system while flushing only gaurentees that data has left Flume's client buffers. When the <code>kite.flushable.commitOnBatch</code> property is set to <code>false</code> , this property must also be set to <code>false</code> . |
| kite.entityParser | avro | Parser that turns Flume Events into Kite entities. Valid values are <code>avro</code> and the fully-qualified class name of an implementation of the <code>EntityParser.Builder</code> interface. |
| kite.failurePolicy | retry | Policy that handles non-recoverable errors such as a missing Schema in the Event header. The default value, <code>retry</code> , will fail the current batch and try again which matches the old behavior. |

| Property Name | Default | Description |
|-------------------------------------|---------|--|
| | | Other valid values are <code>save</code> , which will write the raw Event to the <code>kite.error.dataset.uri</code> dataset, and the fully-qualified class name of an implementation of the <code>FailurePolicy.Builder</code> interface. |
| <code>kite.error.dataset.uri</code> | – | URI of the dataset where failed events are saved when <code>kite.failurePolicy</code> is set to <code>save</code> . Required when the <code>kite.failurePolicy</code> is set to <code>save</code> . |
| <code>auth.kerberosPrincipal</code> | – | Kerberos user principal for secure authentication to HDFS |
| <code>auth.kerberosKeytab</code> | – | Kerberos keytab location (local FS) for the principal |
| <code>auth.proxyUser</code> | – | The effective user for HDFS actions, if different from the kerberos principal |

Kafka Sink

This is a Flume Sink implementation that can publish data to a [Kafka](#) topic. One of the objective is to integrate Flume with Kafka so that pull based processing systems can process the data coming through various Flume sources. This currently supports Kafka 0.8.x series of releases.

Required properties are marked in bold font.

| Property Name | Default | Description |
|-------------------|---------------------|--|
| type | – | Must be set to <code>org.apache.flume.sink.kafka.KafkaSink</code> |
| brokerList | – | List of brokers Kafka-Sink will connect to, to get the list of topic partitions This can be a partial list of brokers, but we recommend at least two for HA. The format is comma separated list of hostname:port |
| topic | default-flume-topic | The topic in Kafka to which the messages will be published. If this parameter is configured, messages will be published to this topic. If the event header |

| Property Name | Default | Description |
|---------------------------------|---------|--|
| | | contains a “topic” field, the event will be published to that topic overriding the topic configured here. |
| batchSize | 100 | How many messages to process in one batch. Larger batches improve throughput while adding latency. |
| requiredAcks | 1 | How many replicas must acknowledge a message before its considered successfully written. Accepted values are 0 (Never wait for acknowledgement), 1 (wait for leader only), -1 (wait for all replicas) Set this to -1 to avoid data loss in some cases of leader failure. |
| Other Kafka Producer Properties | – | These properties are used to configure the Kafka Producer. Any producer property supported by Kafka can be used. The only requirement is to prepend the property name with the prefix <code>kafka.</code> For example: <code>kafka.producer.type</code> |

Note

Kafka Sink uses the `topic` and `key` properties from the FlumeEvent headers to send events to Kafka. If `topic` exists in the headers, the event will be sent to that specific topic, overriding the topic configured for the Sink. If `key` exists in the headers, the key will be used by Kafka to partition the data between the topic partitions. Events with same key will be sent to the same partition. If the key is null, events will be sent to random partitions.

An example configuration of a Kafka sink is given below. Properties starting with the prefix `kafka` (the last 3 properties) are used when instantiating the Kafka producer. The properties that are passed when creating the Kafka producer are not limited to the properties given in this example. Also it's possible to include your custom properties here and access them inside the preprocessor through the Flume Context object passed in as a method argument.

```
a1.sinks.k1.type = org.apache.flume.sink.kafka.KafkaSink
a1.sinks.k1.topic = mytopic
a1.sinks.k1.brokerList = localhost:9092
a1.sinks.k1.requiredAcks = 1
a1.sinks.k1.batchSize = 20
a1.sinks.k1.channel = c1
```

Custom Sink

A custom sink is your own implementation of the Sink interface. A custom sink's class and its dependencies must be included in the agent's classpath when starting the Flume agent. The type of the custom sink is its FQCN. Required properties are in **bold**.

| Property Name | Default | Description |
|----------------|---------|--|
| channel | – | |
| type | – | The component type name, needs to be your FQCN |

Example for agent named a1:

```
a1.channels = c1
a1.sinks = k1
a1.sinks.k1.type = org.example.MySink
a1.sinks.k1.channel = c1
```

Flume Channels

Channels are the repositories where the events are staged on a agent. Source adds the events and Sink removes it.

Memory Channel

The events are stored in an in-memory queue with configurable max size. It's ideal for flows that need higher throughput and are prepared to lose the staged data in the event of a agent failures. Required properties are in **bold**.

| Property Name | Default | Description |
|------------------------------|---------|--|
| type | – | The component type name, needs to be <code>memory</code> |
| capacity | 100 | The maximum number of events stored in the channel |
| transactionCapacity | 100 | The maximum number of events the channel will take from a source or give to a sink per transaction |
| keep-alive | 3 | Timeout in seconds for adding or removing an event |
| byteCapacityBufferPercentage | 20 | Defines the percent of buffer between byteCapacity and the estimated total size of |

| Property Name | Default | Description |
|---------------|-----------------|---|
| | | all events in the channel, to account for data in headers. See below. |
| byteCapacity | see description | Maximum total bytes of memory allowed as a sum of all events in this channel. The implementation only counts the Event body, which is the reason for providing the byteCapacityBufferPercentage configuration parameter as well. Defaults to a computed value equal to 80% of the maximum memory available to the JVM (i.e. 80% of the -Xmx value passed on the command line). Note that if you have multiple memory channels on a single JVM, and they happen to hold the same physical events (i.e. if you are using a replicating channel selector from a single source) then those event sizes may be double-counted for channel byteCapacity purposes. Setting this value to 0 will cause this value to fall back to a hard internal limit of about 200 GB. |

Example for agent named a1:

```
a1.channels = c1
a1.channels.c1.type = memory
a1.channels.c1.capacity = 10000
a1.channels.c1.transactionCapacity = 10000
a1.channels.c1.byteCapacityBufferPercentage = 20
a1.channels.c1.byteCapacity = 800000
```

JDBC Channel

The events are stored in a persistent storage that's backed by a database. The JDBC channel currently supports embedded Derby. This is a durable channel that's ideal for flows where recoverability is important. Required properties are in **bold**.

| Property Name | Default | Description |
|---------------|---------|---|
| type | — | The component type name, needs to be jdbc |

| Property Name | Default | Description |
|----------------------------|--------------------------------------|---|
| db.type | DERBY | Database vendor, needs to be DERBY. |
| driver.class | org.apache.derby.jdbc.EmbeddedDriver | Class for vendor's JDBC driver |
| driver.url | (constructed from other properties) | JDBC connection URL |
| db.username | "sa" | User id for db connection |
| db.password | – | password for db connection |
| connection.properties.file | – | JDBC Connection property file path |
| create.schema | true | If true, then creates db schema if not there |
| create.index | true | Create indexes to speed up lookups |
| create.foreignkey | true | |
| transaction.isolation | "READ_COMMITTED" | Isolation level for db session READ_UNCOMMITTED, READ_COMMITTED, SERIALIZABLE, REPEATABLE_READ |
| maximum.connections | 10 | Max connections allowed to db |
| maximum.capacity | 0 (unlimited) | Max number of events in the channel |
| sysprop.* | | DB Vendor specific properties |
| sysprop.user.home | | Home path to store embedded Derby database |

Example for agent named a1:

```
a1.channels = c1
a1.channels.c1.type = jdbc
```

Kafka Channel

The events are stored in a Kafka cluster (must be installed separately). Kafka provides high availability and replication, so in case an agent or a kafka broker crashes, the events are immediately available to other sinks

The Kafka channel can be used for multiple scenarios: * With Flume source and sink - it provides a reliable and highly available channel for events * With Flume source and interceptor but no sink - it allows writing Flume events into a Kafka topic, for use by other apps * With Flume sink, but no source - it is a low-latency, fault tolerant way to send events from Kafka to Flume sources such as HDFS, HBase or Solr

Required properties are in **bold**.

| Property Name | Default | Description |
|--------------------------|---------------|--|
| type | – | The component type name, needs to be <code>org.apache.flume.channel.kafka.KafkaChannel</code> |
| brokerList | – | List of brokers in the Kafka cluster used by the channel This can be a partial list of brokers, but we recommend at least two for HA. The format is comma separated list of hostname:port |
| zookeeperConnect | – | URI of ZooKeeper used by Kafka cluster The format is comma separated list of hostname:port. If chroot is used, it is added once at the end. For example: <code>zookeeper-1:2181,zookeeper-2:2182,zookeeper-3:2181/kafka</code> |
| topic | flume-channel | Kafka topic which the channel will use |
| groupId | flume | Consumer group ID the channel uses to register with Kafka. Multiple channels must use the same topic and group to ensure that when one agent fails another can get the data Note that having non-channel consumers with the same ID can lead to data loss. |
| parseAsFlumeEvent | true | Expecting Avro datums with FlumeEvent schema in the channel. This should be true if Flume source is writing to the channel And false if other producers are writing into the topic that the channel is using Flume source messages to Kafka can be parsed |

| Property Name | Default | Description |
|------------------------|---------|---|
| | | outside of Flume by using org.apache.flume.source.avro.AvroFlumeEvent provided by the flume-ng-sdk artifact |
| readSmallestOffset | false | When set to true, the channel will read all data in the topic, starting from the oldest event when false, it will read only events written after the channel started When “parseAsFlumeEvent” is true, this will be false. Flume source will start prior to the sinks and this guarantees that events sent by source before sinks start will not be lost. |
| Other Kafka Properties | – | These properties are used to configure the Kafka Producer and Consumer used by the channel. Any property supported by Kafka can be used. The only requirement is to prepend the property name with the prefix <code>kafka.</code> . For example: <code>kafka.producer.type</code> |

Note

Due to the way the channel is load balanced, there may be duplicate events when the agent first starts up

Example for agent named a1:

```
a1.channels.channel1.type =
org.apache.flume.channel.kafka.KafkaChannel
a1.channels.channel1.capacity = 10000
a1.channels.channel1.transactionCapacity = 1000
a1.channels.channel1.brokerList=kafka-2:9092,kafka-3:9092
a1.channels.channel1.topic=channel1
a1.channels.channel1.zookeeperConnect=kafka-1:2181
```

File Channel

Required properties are in **bold**.

| Property Name | Default | Description |
|---------------|---------|--|
| type | – | The component type name, needs to be <code>file</code> . |

| Property Name | Default | Description |
|----------------------|---------------------------------|---|
| checkpointDir | ~/flume/file-channel/checkpoint | The directory where checkpoint file will be stored |
| useDualCheckpoints | false | Backup the checkpoint. If this is set to true, backupCheckpointDir must be set |
| backupCheckpointDir | – | The directory where the checkpoint is backed up to. This directory must not be the same as the data directories or the checkpoint directory |
| dataDirs | ~/flume/file-channel/data | Comma separated list of directories for storing log files. Using multiple directories on separate disks can improve file channel performance |
| transactionCapacity | 10000 | The maximum size of transaction supported by the channel |
| checkpointInterval | 30000 | Amount of time (in millis) between checkpoints |
| maxFileSize | 2146435071 | Max size (in bytes) of a single log file |
| minimumRequiredSpace | 524288000 | Minimum Required free space (in bytes). To avoid data corruption, File Channel stops accepting take/put requests when free space drops below this value |
| capacity | 1000000 | Maximum capacity of the channel |
| keep-alive | 3 | Amount of time (in sec) to wait for a put operation |
| use-log-replay-v1 | false | Expert: Use old replay logic |
| use-fast-replay | false | Expert: Replay without using queue |
| checkpointOnClose | true | Controls if a checkpoint is created when the channel is closed. Creating a checkpoint on close speeds up subsequent |

| Property Name | Default | Description |
|---|---------|--|
| | | startup of the file channel by avoiding replay. |
| encryption.activeKey | — | Key name used to encrypt new data |
| encryption.cipherProvider | — | Cipher provider type, supported types: AESCTRNOPADDING |
| encryption.keyProvider | — | Key provider type, supported types: JCEKSFILE |
| encryption.keyProvider.keyStoreFile | — | Path to the keystore file |
| encryption.keyProvider.keyStorePasswordFile | — | Path to the keystore password file |
| encryption.keyProvider.keys | — | List of all keys (e.g. history of the activeKey setting) |
| encryption.keyProvider.keys.*.passwordFile | — | Path to the optional key password file |

Note

By default the File Channel uses paths for checkpoint and data directories that are within the user home as specified above. As a result if you have more than one File Channel instances active within the agent, only one will be able to lock the directories and cause the other channel initialization to fail. It is therefore necessary that you provide explicit paths to all the configured channels, preferably on different disks. Furthermore, as file channel will sync to disk after every commit, coupling it with a sink/source that batches events together may be necessary to provide good performance where multiple disks are not available for checkpoint and data directories.

Example for agent named a1:

```
a1.channels = c1
a1.channels.c1.type = file
a1.channels.c1.checkpointDir = /mnt/flume/checkpoint
a1.channels.c1.dataDirs = /mnt/flume/data
```

Encryption

Below is a few sample configurations:

Generating a key with a password separate from the key store password:

```
keytool -genseckey -alias key-0 -keypass keyPassword -keyalg AES \
  -keysize 128 -validity 9000 -keystore test.keystore \
  -storetype jceks -storepass keyStorePassword
```

Generating a key with the password the same as the key store password:

```
keytool -genseckey -alias key-1 -keyalg AES -keysize 128 -validity
9000 \
  -keystore src/test/resources/test.keystore -storetype jceks \
  -storepass keyStorePassword
a1.channels.c1.encryption.activeKey = key-0
a1.channels.c1.encryption.cipherProvider = AESCTRNOPADDING
a1.channels.c1.encryption.keyProvider = key-provider-0
a1.channels.c1.encryption.keyProvider = JCEKSFILE
a1.channels.c1.encryption.keyProvider.keyStoreFile =
/path/to/my.keystore
a1.channels.c1.encryption.keyProvider.keyStorePasswordFile =
/path/to/my.keystore.password
a1.channels.c1.encryption.keyProvider.keys = key-0
```

Let's say you have aged key-0 out and new files should be encrypted with key-1:

```
a1.channels.c1.encryption.activeKey = key-1
a1.channels.c1.encryption.cipherProvider = AESCTRNOPADDING
a1.channels.c1.encryption.keyProvider = JCEKSFILE
a1.channels.c1.encryption.keyProvider.keyStoreFile =
/path/to/my.keystore
a1.channels.c1.encryption.keyProvider.keyStorePasswordFile =
/path/to/my.keystore.password
a1.channels.c1.encryption.keyProvider.keys = key-0 key-1
```

The same scenario as above, however key-0 has its own password:

```
a1.channels.c1.encryption.activeKey = key-1
a1.channels.c1.encryption.cipherProvider = AESCTRNOPADDING
a1.channels.c1.encryption.keyProvider = JCEKSFILE
a1.channels.c1.encryption.keyProvider.keyStoreFile =
/path/to/my.keystore
a1.channels.c1.encryption.keyProvider.keyStorePasswordFile =
/path/to/my.keystore.password
a1.channels.c1.encryption.keyProvider.keys = key-0 key-1
a1.channels.c1.encryption.keyProvider.keys.key-0.passwordFile =
/path/to/key-0.password
```

Spillable Memory Channel

The events are stored in an in-memory queue and on disk. The in-memory queue serves as the primary store and the disk as overflow. The disk store is managed using an embedded File channel. When the in-memory queue is full, additional incoming events are stored in the file channel. This channel is ideal for flows that need high throughput of memory channel during normal operation, but at the same time need the larger capacity of the file channel for better tolerance of intermittent sink side outages or drop in drain rates. The throughput will reduce approximately to file channel speeds during such abnormal situations. In case of an agent crash or restart, only the events stored on disk are recovered when the agent comes online. **This channel is currently experimental and not recommended for use in production.**

Required properties are in **bold**. Please refer to file channel for additional required properties.

| Property Name | Default | Description |
|------------------------------|-----------------|---|
| type | – | The component type name, needs to be <code>SPILLABLEMEMORY</code> |
| memoryCapacity | 10000 | Maximum number of events stored in memory queue. To disable use of in-memory queue, set this to zero. |
| overflowCapacity | 1000000 | Maximum number of events stored in overflow disk (i.e File channel). To disable use of overflow, set this to zero. |
| overflowTimeout | 3 | The number of seconds to wait before enabling disk overflow when memory fills up. |
| byteCapacityBufferPercentage | 20 | Defines the percent of buffer between byteCapacity and the estimated total size of all events in the channel, to account for data in headers. See below. |
| byteCapacity | see description | Maximum bytes of memory allowed as a sum of all events in the memory queue. The implementation only counts the Event body , which is the reason for providing the <code>byteCapacityBufferPercentage</code> configuration parameter as well. Defaults to a computed value equal to 80% of the maximum memory available to the JVM |

| Property Name | Default | Description |
|---------------------------|------------------|--|
| | | (i.e. 80% of the -Xmx value passed on the command line). Note that if you have multiple memory channels on a single JVM, and they happen to hold the same physical events (i.e. if you are using a replicating channel selector from a single source) then those event sizes may be double-counted for channel byteCapacity purposes. Setting this value to 0 will cause this value to fall back to a hard internal limit of about 200 GB. |
| avgEventSize | 500 | Estimated average size of events, in bytes, going into the channel |
| <file channel properties> | see file channel | Any file channel property with the exception of 'keep-alive' and 'capacity' can be used. The keep-alive of file channel is managed by Spillable Memory Channel. Use 'overflowCapacity' to set the File channel's capacity. |

In-memory queue is considered full if either memoryCapacity or byteCapacity limit is reached.

Example for agent named a1:

```
a1.channels = c1
a1.channels.c1.type = SPILLABLEMEMORY
a1.channels.c1.memoryCapacity = 10000
a1.channels.c1.overflowCapacity = 1000000
a1.channels.c1.byteCapacity = 800000
a1.channels.c1.checkpointDir = /mnt/flume/checkpoint
a1.channels.c1.dataDirs = /mnt/flume/data
```

To disable the use of the in-memory queue and function like a file channel:

```
a1.channels = c1
a1.channels.c1.type = SPILLABLEMEMORY
a1.channels.c1.memoryCapacity = 0
a1.channels.c1.overflowCapacity = 1000000
a1.channels.c1.checkpointDir = /mnt/flume/checkpoint
a1.channels.c1.dataDirs = /mnt/flume/data
```

To disable the use of overflow disk and function purely as a in-memory channel:

```
a1.channels = c1
a1.channels.c1.type = SPILLABLEMEMORY
a1.channels.c1.memoryCapacity = 100000
a1.channels.c1.overflowCapacity = 0
```

Pseudo Transaction Channel

Warning

The Pseudo Transaction Channel is only for unit testing purposes and is NOT meant for production use.

Required properties are in **bold**.

| Property Name | Default | Description |
|---------------|---------|---|
| type | – | The component type name, needs to be <code>org.apache.flume.channel.PseudoTxnMemoryChannel</code> |
| capacity | 50 | The max number of events stored in the channel |
| keep-alive | 3 | Timeout in seconds for adding or removing an event |

Custom Channel

A custom channel is your own implementation of the Channel interface. A custom channel's class and its dependencies must be included in the agent's classpath when starting the Flume agent. The type of the custom channel is its FQCN. Required properties are in **bold**.

| Property Name | Default | Description |
|---------------|---------|---|
| type | – | The component type name, needs to be a FQCN |

Example for agent named a1:

```
a1.channels = c1
a1.channels.c1.type = org.example.MyChannel
```

Flume Channel Selectors

If the type is not specified, then defaults to “replicating”.

Replicating Channel Selector (default)

Required properties are in **bold**.

| Property Name | Default | Description |
|-------------------|-------------|--|
| selector.type | replicating | The component type name, needs to be replicating |
| selector.optional | – | Set of channels to be marked as optional |

Example for agent named a1 and it's source called r1:

```
a1.sources = r1
a1.channels = c1 c2 c3
a1.source.r1.selector.type = replicating
a1.source.r1.channels = c1 c2 c3
a1.source.r1.selector.optional = c3
```

In the above configuration, c3 is an optional channel. Failure to write to c3 is simply ignored. Since c1 and c2 are not marked optional, failure to write to those channels will cause the transaction to fail.

Multiplexing Channel Selector

Required properties are in **bold**.

| Property Name | Default | Description |
|--------------------|-----------------------|---|
| selector.type | replicating | The component type name, needs to be multiplexing |
| selector.header | flume.selector.header | |
| selector.default | – | |
| selector.mapping.* | – | |

Example for agent named a1 and it's source called r1:

```
a1.sources = r1
a1.channels = c1 c2 c3 c4
a1.sources.r1.selector.type = multiplexing
a1.sources.r1.selector.header = state
a1.sources.r1.selector.mapping.CZ = c1
```

```
a1.sources.r1.selector.mapping.US = c2 c3  
a1.sources.r1.selector.default = c4
```

Custom Channel Selector

A custom channel selector is your own implementation of the ChannelSelector interface. A custom channel selector's class and its dependencies must be included in the agent's classpath when starting the Flume agent. The type of the custom channel selector is its FQCN.

| Property Name | Default | Description |
|---------------|---------|--|
| selector.type | – | The component type name, needs to be your FQCN |

Example for agent named a1 and its source called r1:

```
a1.sources = r1  
a1.channels = c1  
a1.sources.r1.selector.type = org.example.MyChannelSelector
```

Flume Sink Processors

Sink groups allow users to group multiple sinks into one entity. Sink processors can be used to provide load balancing capabilities over all sinks inside the group or to achieve fail over from one sink to another in case of temporal failure.

Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------------|---------|--|
| sinks | – | Space-separated list of sinks that are participating in the group |
| processor.type | default | The component type name, needs to be default, failover or load_balance |

Example for agent named a1:

```
a1.sinkgroups = g1  
a1.sinkgroups.g1.sinks = k1 k2  
a1.sinkgroups.g1.processor.type = load_balance
```


Default Sink Processor

Default sink processor accepts only a single sink. User is not forced to create processor (sink group) for single sinks. Instead user can follow the source - channel - sink pattern that was explained above in this user guide.

Failover Sink Processor

Failover Sink Processor maintains a prioritized list of sinks, guaranteeing that so long as one is available events will be processed (delivered).

The failover mechanism works by relegating failed sinks to a pool where they are assigned a cool down period, increasing with sequential failures before they are retried. Once a sink successfully sends an event, it is restored to the live pool. The Sinks have a priority associated with them, larger the number, higher the priority. If a Sink fails while sending a Event the next Sink with highest priority shall be tried next for sending Events. For example, a sink with priority 100 is activated before the Sink with priority 80. If no priority is specified, the priority is determined based on the order in which the Sinks are specified in configuration.

To configure, set a sink groups processor to `failover` and set priorities for all individual sinks. All specified priorities must be unique. Furthermore, upper limit to failover time can be set (in milliseconds) using `maxpenalty` property.

Required properties are in **bold**.

| Property Name | Default | Description |
|--|----------------------|--|
| sinks | – | Space-separated list of sinks that are participating in the group |
| processor.type | <code>default</code> | The component type name, needs to be <code>failover</code> |
| processor.priority.<sinkName> | – | Priority value. <sinkName> must be one of the sink instances associated with the current sink group A higher priority value Sink gets activated earlier. A larger absolute value indicates higher priority |
| processor.maxpenalty | 30000 | The maximum backoff period for the failed Sink (in millis) |

Example for agent named a1:

```
a1.sinkgroups = g1
a1.sinkgroups.g1.sinks = k1 k2
a1.sinkgroups.g1.processor.type = failover
a1.sinkgroups.g1.processor.priority.k1 = 5
a1.sinkgroups.g1.processor.priority.k2 = 10
a1.sinkgroups.g1.processor.maxpenalty = 10000
```

Load balancing Sink Processor

Load balancing sink processor provides the ability to load-balance flow over multiple sinks. It maintains an indexed list of active sinks on which the load must be distributed. Implementation supports distributing load using either via `round_robin` or `random` selection mechanisms. The choice of selection mechanism defaults to `round_robin` type, but can be overridden via configuration. Custom selection mechanisms are supported via custom classes that inherits from `AbstractSinkSelector`.

When invoked, this selector picks the next sink using its configured selection mechanism and invokes it. For `round_robin` and `random` In case the selected sink fails to deliver the event, the processor picks the next available sink via its configured selection mechanism. This implementation does not blacklist the failing sink and instead continues to optimistically attempt every available sink. If all sinks invocations result in failure, the selector propagates the failure to the sink runner.

If `backoff` is enabled, the sink processor will blacklist sinks that fail, removing them for selection for a given timeout. When the timeout ends, if the sink is still unresponsive timeout is increased exponentially to avoid potentially getting stuck in long waits on unresponsive sinks. With this disabled, in round-robin all the failed sinks load will be passed to the next sink in line and thus not evenly balanced

Required properties are in **bold**.

| Property Name | Default | Description |
|------------------------|---------|---|
| processor.sinks | – | Space-separated list of sinks that are participating in the group |
| processor.type | default | The component type name, needs to be <code>load_balance</code> |

| Property Name | Default | Description |
|-------------------------------|-------------|---|
| processor.backoff | false | Should failed sinks be backed off exponentially. |
| processor.selector | round_robin | Selection mechanism. Must be either round_robin, random or FQCN of custom class that inherits from AbstractSinkSelector |
| processor.selector.maxTimeOut | 30000 | Used by backoff selectors to limit exponential backoff (in milliseconds) |

Example for agent named a1:

```
a1.sinkgroups = g1
a1.sinkgroups.g1.sinks = k1 k2
a1.sinkgroups.g1.processor.type = load_balance
a1.sinkgroups.g1.processor.backoff = true
a1.sinkgroups.g1.processor.selector = random
```

Custom Sink Processor

Custom sink processors are not supported at the moment.

Event Serializers

The file_roll sink and the hdfs sink both support the EventSerializer interface. Details of the EventSerializers that ship with Flume are provided below.

Body Text Serializer

Alias: text. This interceptor writes the body of the event to an output stream without any transformation or modification. The event headers are ignored. Configuration options are as follows:

| Property Name | Default | Description |
|---------------|---------|--|
| appendNewline | true | Whether a newline will be appended to each event at write time. The default of true assumes that events do not contain newlines, for legacy reasons. |

Example for agent named a1:

```
a1.sinks.k1 = k1
a1.sinks.k1.type = file_roll
a1.sinks.k1.channel = c1
a1.sinks.k1.sink.directory = /var/log/flume
a1.sinks.k1.sink.serializer = text
a1.sinks.k1.sink.serializer.appendNewline = false
```

Avro Event Serializer

Alias: `avro_event`. This interceptor serializes Flume events into an Avro container file. The schema used is the same schema used for Flume events in the Avro RPC mechanism. This serializers inherits from the `AbstractAvroEventSerializer` class. Configuration options are as follows:

| Property Name | Default | Description |
|-------------------|---------|---|
| syncIntervalBytes | 2048000 | Avro sync interval, in approximate bytes. |
| compressionCodec | null | Avro compression codec. For supported codecs, see Avro's CodecFactory docs. |

Example for agent named a1:

```
a1.sinks.k1.type = hdfs
a1.sinks.k1.channel = c1
a1.sinks.k1.hdfs.path = /flume/events/%y-%m-%d/%H%M/%S
a1.sinks.k1.serializer = avro_event
a1.sinks.k1.serializer.compressionCodec = snappy
```

Flume Interceptors

Flume has the capability to modify/drop events in-flight. This is done with the help of interceptors. Interceptors are classes that implement `org.apache.flume.interceptor.Interceptor` interface. An

interceptor can modify or even drop events based on any criteria chosen by the developer of the interceptor. Flume supports chaining of interceptors. This is made possible through by specifying the list of interceptor builder class names in the configuration. Interceptors are specified as a whitespace separated list in the source configuration. The order in which the interceptors are specified is the order in which they are invoked. The list of events returned by one interceptor is passed to the next interceptor in the chain. Interceptors can modify or drop events. If an interceptor needs to drop events, it just does not return that event in the list that it returns. If it is to drop all events, then it simply returns an empty list. Interceptors are named components, here is an example of how they are created through configuration:

```
al.sources = r1
al.sinks = k1
al.channels = c1
al.sources.r1.interceptors = i1 i2
al.sources.r1.interceptors.i1.type =
org.apache.flume.interceptor.HostInterceptor$Builder
al.sources.r1.interceptors.i1.preserveExisting = false
al.sources.r1.interceptors.i1.hostHeader = hostname
al.sources.r1.interceptors.i2.type =
org.apache.flume.interceptor.TimestampInterceptor$Builder
al.sinks.k1.filePrefix = FlumeData.{CollectorHost}.%Y-%m-%d
al.sinks.k1.channel = c1
```

Note that the interceptor builders are passed to the type config parameter. The interceptors are themselves configurable and can be passed configuration values just like they are passed to any other configurable component. In the above example, events are passed to the HostInterceptor first and the events returned by the HostInterceptor are then passed along to the TimestampInterceptor. You can specify either the fully qualified class name (FQCN) or the alias `timestamp`. If you have multiple collectors writing to the same HDFS path, then you could also use the HostInterceptor.

Timestamp Interceptor

This interceptor inserts into the event headers, the time in millis at which it processes the event. This interceptor inserts a header with key `timestamp` whose value is the relevant timestamp. This interceptor can preserve an existing timestamp if it is already present in the configuration.

| Property Name | Default | Description |
|------------------|---------|---|
| type | – | The component type name, has to be <code>timestamp</code> or the FQCN |
| preserveExisting | false | If the timestamp already exists, should it be preserved - true or false |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.channels = c1
a1.sources.r1.type = seq
a1.sources.r1.interceptors = i1
a1.sources.r1.interceptors.i1.type = timestamp
```

Host Interceptor

This interceptor inserts the hostname or IP address of the host that this agent is running on. It inserts a header with key `host` or a configured key whose value is the hostname or IP address of the host, based on configuration.

| Property Name | Default | Description |
|------------------|---------|---|
| type | – | The component type name, has to be <code>host</code> |
| preserveExisting | false | If the host header already exists, should it be preserved - true or false |
| useIP | true | Use the IP Address if true, else use hostname. |
| hostHeader | host | The header key to be used. |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.interceptors = i1
a1.sources.r1.interceptors.i1.type = host
a1.sources.r1.interceptors.i1.hostHeader = hostname
```

Static Interceptor

Static interceptor allows user to append a static header with static value to all events.

The current implementation does not allow specifying multiple headers at one time. Instead user might chain multiple static interceptors each defining one static header.

| Property Name | Default | Description |
|------------------|---------|---|
| type | – | The component type name, has to be <code>static</code> |
| preserveExisting | true | If configured header already exists, should it be preserved - true or false |
| key | key | Name of header that should be created |
| value | value | Static value that should be created |

Example for agent named a1:

```
a1.sources = r1
a1.channels = c1
a1.sources.r1.channels = c1
a1.sources.r1.type = seq
a1.sources.r1.interceptors = i1
a1.sources.r1.interceptors.i1.type = static
a1.sources.r1.interceptors.i1.key = datacenter
a1.sources.r1.interceptors.i1.value = NEW_YORK
```

UUID Interceptor

This interceptor sets a universally unique identifier on all events that are intercepted. An example UUID is `b5755073-77a9-43c1-8fad-b7a586fc1b97`, which represents a 128-bit value.

Consider using `UUIDInterceptor` to automatically assign a UUID to an event if no application level unique key for the event is available. It can be important to assign UUIDs to events as soon as they enter the Flume network; that is, in the first Flume Source of the flow. This enables subsequent deduplication of events in the face of replication and redelivery in a Flume network that is designed for high availability and high performance. If an application level key is available, this is preferable over an auto-generated UUID because it enables subsequent updates and deletes of event in data stores using said well known application level key.

| Property Name | Default | Description |
|---------------|---------|--|
| type | – | The component type name has to be <code>org.apache.flume.sink.solr.morphline.UUIDInterceptor\$Builder</code> |

| Property Name | Default | Description |
|------------------|---------|---|
| headerName | id | The name of the Flume header to modify |
| preserveExisting | true | If the UUID header already exists, should it be preserved - true or false |
| prefix | "" | The prefix string constant to prepend to each generated UUID |

Morphline Interceptor

This interceptor filters the events through a [morphline configuration file](#) that defines a chain of transformation commands that pipe records from one command to another. For example the morphline can ignore certain events or alter or insert certain event headers via regular expression based pattern matching, or it can auto-detect and set a MIME type via Apache Tika on events that are intercepted. For example, this kind of packet sniffing can be used for content based dynamic routing in a Flume topology. MorphlineInterceptor can also help to implement dynamic routing to multiple Apache Solr collections (e.g. for multi-tenancy).

Currently, there is a restriction in that the morphline of an interceptor must not generate more than one output record for each input event. This interceptor is not intended for heavy duty ETL processing - if you need this consider moving ETL processing from the Flume Source to a Flume Sink, e.g. to a MorphlineSolrSink.

Required properties are in **bold**.

| Property Name | Default | Description |
|----------------------|---------|---|
| type | – | The component type name has to be <code>org.apache.flume.sink.solr.morphline.MorphlineInterceptor\$Builder</code> |
| morphlineFile | – | The relative or absolute path on the local file system to the morphline configuration file. Example: <code>/etc/flume-ng/conf/morphline.conf</code> |
| morphlineId | null | Optional name used to identify a morphline if there are multiple morphlines in a morphline config file |

Sample flume.conf file:


```
al.sources.avroSrc.interceptors = morphlineinterceptor
al.sources.avroSrc.interceptors.morphlineinterceptor.type =
org.apache.flume.sink.solr.morphline.MorphlineInterceptor$Builder
al.sources.avroSrc.interceptors.morphlineinterceptor.morphlineFile
= /etc/flume-ng/conf/morphline.conf
al.sources.avroSrc.interceptors.morphlineinterceptor.morphlineId =
morphline1
```

Search and Replace Interceptor

This interceptor provides simple string-based search-and-replace functionality based on Java regular expressions. Backtracking / group capture is also available. This interceptor uses the same rules as in the Java `Matcher.replaceAll()` method.

| Property Name | Default | Description |
|----------------------------|---------|--|
| type | – | The component type name has to be <code>search_replace</code> |
| <code>searchPattern</code> | – | The pattern to search for and replace. |
| <code>replaceString</code> | – | The replacement string. |
| <code>charset</code> | UTF-8 | The charset of the event body. Assumed by default to be UTF-8. |

Example configuration:

```
al.sources.avroSrc.interceptors = search-replace
al.sources.avroSrc.interceptors.search-replace.type =
search_replace

# Remove leading alphanumeric characters in an event body.
al.sources.avroSrc.interceptors.search-replace.searchPattern =
^[A-Za-z0-9_]+
al.sources.avroSrc.interceptors.search-replace.replaceString =
```

Another example:

```
al.sources.avroSrc.interceptors = search-replace
al.sources.avroSrc.interceptors.search-replace.type =
search_replace

# Use grouping operators to reorder and munge words on a line.
al.sources.avroSrc.interceptors.search-replace.searchPattern = The
quick brown ([a-z]+) jumped over the lazy ([a-z]+)
```

```
al.sources.avroSrc.interceptors.search-replace.replaceString = The
hungry $2 ate the careless $1
```

Regex Filtering Interceptor

This interceptor filters events selectively by interpreting the event body as text and matching the text against a configured regular expression. The supplied regular expression can be used to include events or exclude events.

| Property Name | Default | Description |
|---------------|---------------------|--|
| type | – | The component type name has to be <code>regex_filter</code> |
| regex | " <code>.*</code> " | Regular expression for matching against events |
| excludeEvents | false | If true, regex determines events to exclude, otherwise regex determines events to include. |

Regex Extractor Interceptor

This interceptor extracts regex match groups using a specified regular expression and appends the match groups as headers on the event. It also supports pluggable serializers for formatting the match groups before adding them as event headers.

| Property Name | Default | Description |
|-----------------------|---------|--|
| type | – | The component type name has to be <code>regex_extractor</code> |
| regex | – | Regular expression for matching against events |
| serializers | – | Space-separated list of serializers for mapping matches to header names and serializing their values. (See example below) Flume provides built-in support for the following serializers: <code>org.apache.flume.interceptor.RegexExtractorInterceptorPassThroughSerializer</code> <code>org.apache.flume.interceptor.RegexExtractorInterceptorMillisSerializer</code> |
| serializers.<sl>.type | default | Must be default (<code>org.apache.flume.interceptor.RegexExtractorInterceptorPassThroughSerializer</code>), <code>org.apache.flume.interceptor.RegexExtractorInterceptorMillisSerializer</code> , or the |

| Property Name | Default | Description |
|--|---------|--|
| | | FQCN of a custom class that implements <code>org.apache.flume.interceptor.RegexExtractorInterceptorSerializer</code> |
| <code>serializers.<s1>.name</code> | – | |
| <code>serializers.*</code> | – | Serializer-specific properties |

The serializers are used to map the matches to a header name and a formatted header value; by default, you only need to specify the header name and the default `org.apache.flume.interceptor.RegexExtractorInterceptorPassThroughSerializer` will be used. This serializer simply maps the matches to the specified header name and passes the value through as it was extracted by the regex. You can plug custom serializer implementations into the extractor using the fully qualified class name (FQCN) to format the matches in anyway you like.

Example 1:

If the Flume event body contained `1:2:3.4foobar5` and the following configuration was used

```
a1.sources.r1.interceptors.i1.regex = (\\d):(\\d):(\\d)
a1.sources.r1.interceptors.i1.serializers = s1 s2 s3
a1.sources.r1.interceptors.i1.serializers.s1.name = one
a1.sources.r1.interceptors.i1.serializers.s2.name = two
a1.sources.r1.interceptors.i1.serializers.s3.name = three
```

The extracted event will contain the same body but the following headers will have been added `one=>1, two=>2, three=>3`

Example 2:

If the Flume event body contained 2012-10-18 18:47:57,614 some log line and the following configuration was used

```
a1.sources.r1.interceptors.i1.regex = ^(?:\\n)?(\\d\\d\\d\\d-\\d\\d\\d-\\d\\d\\d\\s\\d\\d\\d:\\d\\d)
a1.sources.r1.interceptors.i1.serializers = s1
a1.sources.r1.interceptors.i1.serializers.s1.type =
org.apache.flume.interceptor.RegexExtractorInterceptorMillisSerial
izer
a1.sources.r1.interceptors.i1.serializers.s1.name = timestamp
a1.sources.r1.interceptors.i1.serializers.s1.pattern = yyyy-MM-dd
HH:mm
```

the extracted event will contain the same body but the following headers will have been added timestamp=>1350611220000

Flume Properties

| Property Name | Default | Description |
|---------------------------|---------|---|
| flume.called.from.service | - | If this property is specified then the Flume agent will continue polling for the config file even if the config file is not found at the expected location. Otherwise, the Flume agent will terminate if the config doesn't exist at the expected location. No property value is needed when setting this property (eg, just specifying - Dflume.called.from.service is enough) |

Property: flume.called.from.service

Flume periodically polls, every 30 seconds, for changes to the specified config file. A Flume agent loads a new configuration from the config file if either an existing file is polled for the first time, or if an existing file's modification date has changed since the last time it was polled. Renaming or moving a file does not change its modification time. When a Flume agent polls a non-existent file then one of two things happens: 1. When the agent polls a non-existent config file for the first time, then the agent behaves according to the flume.called.from.service property. If the property is set, then the agent will continue polling (always at the same period – every 30 seconds). If the property is not set, then the agent immediately terminates. ...OR... 2. When the agent polls a non-existent config file and this is not the first time

the file is polled, then the agent makes no config changes for this polling period. The agent continues polling rather than terminating.

Log4J Appender

Appends Log4j events to a flume agent's avro source. A client using this appender must have the flume-ng-sdk in the classpath (eg, flume-ng-sdk-1.7.0-SNAPSHOT.jar). Required properties are in **bold**.

| Property Name | Default | Description |
|-----------------------|---------|--|
| Hostname | – | The hostname on which a remote Flume agent is running with an avro source. |
| Port | – | The port at which the remote Flume agent's avro source is listening. |
| UnsafeMode | false | If true, the appender will not throw exceptions on failure to send the events. |
| AvroReflectionEnabled | false | Use Avro Reflection to serialize Log4j events. (Do not use when users log strings) |
| AvroSchemaUrl | – | A URL from which the Avro schema can be retrieved. |

Sample log4j.properties file:

```
#...
log4j.appender.flume =
org.apache.flume.clients.log4jappender.Log4jAppender
log4j.appender.flume.Hostname = example.com
log4j.appender.flume.Port = 41414
log4j.appender.flume.UnsafeMode = true

# configure a class's logger to output to the flume appender
log4j.logger.org.example.MyClass = DEBUG,flume
#...
```

By default each event is converted to a string by calling `toString()`, or by using the Log4j layout, if specified.

If the event is an instance of `org.apache.avro.generic.GenericRecord`, `org.apache.avro.specific.SpecificRecord`, or if the property `AvroReflectionEnabled` is set to `true` then the event will be serialized using Avro serialization.

Serializing every event with its Avro schema is inefficient, so it is good practice to provide a schema URL from which the schema can be retrieved by the downstream sink, typically the HDFS sink. If `AvroSchemaUrl` is not specified, then the schema will be included as a Flume header.

Sample `log4j.properties` file configured to use Avro serialization:

```
#...
log4j.appender.flume =
org.apache.flume.clients.log4jappender.Log4jAppender
log4j.appender.flume.Hostname = example.com
log4j.appender.flume.Port = 41414
log4j.appender.flume.AvroReflectionEnabled = true
log4j.appender.flume.AvroSchemaUrl =
hdfs://namenode/path/to/schema.avsc

# configure a class's logger to output to the flume appender
log4j.logger.org.example.MyClass = DEBUG,flume
#...
```

Load Balancing Log4J Appender

Appends Log4j events to a list of flume agent's avro source. A client using this appender must have the `flume-ng-sdk` in the classpath (eg, `flume-ng-sdk-1.7.0-SNAPSHOT.jar`). This appender supports a round-robin and random scheme for performing the load balancing. It also supports a configurable backoff timeout so that down agents are removed temporarily from the set of hosts Required properties are in **bold**.

| Property Name | Default | Description |
|---------------|-------------|--|
| Hosts | — | A space-separated list of host:port at which Flume (through an AvroSource) is listening for events |
| Selector | ROUND_ROBIN | Selection mechanism. Must be either ROUND_ROBIN, RANDOM or custom FQDN to class that inherits from LoadBalancingSelector. |
| MaxBackoff | — | A long value representing the maximum amount of time in milliseconds the Load balancing client will backoff from a node that |

| Property Name | Default | Description |
|-----------------------|---------|--|
| | | has failed to consume an event. Defaults to no backoff |
| UnsafeMode | false | If true, the appender will not throw exceptions on failure to send the events. |
| AvroReflectionEnabled | false | Use Avro Reflection to serialize Log4j events. |
| AvroSchemaUrl | – | A URL from which the Avro schema can be retrieved. |

Sample log4j.properties file configured using defaults:

```
#...
log4j.appender.out2 =
org.apache.flume.clients.log4jappender.LoadBalancingLog4jAppender
log4j.appender.out2.Hosts = localhost:25430 localhost:25431

# configure a class's logger to output to the flume appender
log4j.logger.org.example.MyClass = DEBUG,flume
#...
```

Sample log4j.properties file configured using RANDOM load balancing:

```
#...
log4j.appender.out2 =
org.apache.flume.clients.log4jappender.LoadBalancingLog4jAppender
log4j.appender.out2.Hosts = localhost:25430 localhost:25431
log4j.appender.out2.Selector = RANDOM

# configure a class's logger to output to the flume appender
log4j.logger.org.example.MyClass = DEBUG,flume
#...
```

Sample log4j.properties file configured using backoff:

```
#...
log4j.appender.out2 =
org.apache.flume.clients.log4jappender.LoadBalancingLog4jAppender
log4j.appender.out2.Hosts = localhost:25430 localhost:25431
localhost:25432
log4j.appender.out2.Selector = ROUND_ROBIN
log4j.appender.out2.MaxBackoff = 30000

# configure a class's logger to output to the flume appender
```

```
log4j.logger.org.example.MyClass = DEBUG, flume  
#...
```

Security

The HDFS sink, HBase sink, Thrift source, Thrift sink and Kite Dataset sink all support Kerberos authentication. Please refer to the corresponding sections for configuring the Kerberos-related options.

Flume agent will authenticate to the kerberos KDC as a single principal, which will be used by different components that require kerberos authentication. The principal and keytab configured for Thrift source, Thrift sink, HDFS sink, HBase sink and DataSet sink should be the same, otherwise the component will fail to start.

Monitoring

Monitoring in Flume is still a work in progress. Changes can happen very often. Several Flume components report metrics to the JMX platform MBean server. These metrics can be queried using Jconsole.

JMX Reporting

JMX Reporting can be enabled by specifying JMX parameters in the JAVA_OPTS environment variable using flume-env.sh, like

```
export JAVA_OPTS="-Dcom.sun.management.jmxremote -  
Dcom.sun.management.jmxremote.port=5445 -  
Dcom.sun.management.jmxremote.authenticate=false -  
Dcom.sun.management.jmxremote.ssl=false"
```

NOTE: The sample above disables the security. To enable Security, please refer <http://docs.oracle.com/javase/6/docs/technotes/guides/management/agent.htm>

1

Ganglia Reporting

Flume can also report these metrics to Ganglia 3 or Ganglia 3.1 metanodes. To report metrics to Ganglia, a flume agent must be started with this support. The Flume agent

has to be started by passing in the following parameters as system properties prefixed by `flume.monitoring.`, and can be specified in the `flume-env.sh`:

| Property Name | Default | Description |
|----------------------|---------|--|
| type | – | The component type name, has to be <code>ganglia</code> |
| hosts | – | Comma-separated list of <code>hostname:port</code> of Ganglia servers |
| pollFrequency | 60 | Time, in seconds, between consecutive reporting to Ganglia server |
| isGanglia3 | false | Ganglia server version is 3. By default, Flume sends in Ganglia 3.1 format |

We can start Flume with Ganglia support as follows:

```
$ bin/flume-ng agent --conf-file example.conf --name a1 -
Dflume.monitoring.type=ganglia -
Dflume.monitoring.hosts=com.example:1234,com.example2:5455
```

JSON Reporting

Flume can also report metrics in a JSON format. To enable reporting in JSON format, Flume hosts a Web server on a configurable port. Flume reports metrics in the following JSON format:

```
{
  "typeName1.componentName1" : {"metric1" : "metricValue1",
  "metric2" : "metricValue2"},
  "typeName2.componentName2" : {"metric3" : "metricValue3",
  "metric4" : "metricValue4"}
}
```

Here is an example:

```
{
  "CHANNEL.fileChannel": {"EventPutSuccessCount": "468085",
    "Type": "CHANNEL",
    "StopTime": "0",
    "EventPutAttemptCount": "468086",
    "ChannelSize": "233428",
    "StartTime": "1344882233070",
    "EventTakeSuccessCount": "458200",
    "ChannelCapacity": "600000",
```

```
        "EventTakeAttemptCount": "458288"},  
"CHANNEL.memChannel": {"EventPutSuccessCount": "22948908",  
    "Type": "CHANNEL",  
    "StopTime": "0",  
    "EventPutAttemptCount": "22948908",  
    "ChannelSize": "5",  
    "StartTime": "1344882209413",  
    "EventTakeSuccessCount": "22948900",  
    "ChannelCapacity": "100",  
    "EventTakeAttemptCount": "22948908"}  
}
```

| Property Name | Default | Description |
|---------------|---------|---|
| type | – | The component type name, has to be http |
| port | 41414 | The port to start the server on. |

We can start Flume with JSON Reporting support as follows:

```
$ bin/flume-ng agent --conf-file example.conf --name a1 -  
Dflume.monitoring.type=http -Dflume.monitoring.port=34545
```

Metrics will then be available at **http://<hostname>:<port>/metrics** webpage.
Custom components can report metrics as mentioned in the Ganglia section above.

Custom Reporting

It is possible to report metrics to other systems by writing servers that do the reporting. Any reporting class has to implement the interface, `org.apache.flume.instrumentation.MonitorService`. Such a class can be used the same way the `GangliaServer` is used for reporting. They can poll the platform mbean server to poll the mbeans for metrics. For example, if an HTTP monitoring service called `HTTPReporting` can be used as follows:

```
$ bin/flume-ng agent --conf-file example.conf --name a1 -  
Dflume.monitoring.type=com.example.reporting.HTTPReporting -  
Dflume.monitoring.node=com.example:332
```

| Property Name | Default | Description |
|---------------|---------|---|
| type | – | The component type name, has to be FQCN |

Reporting metrics from custom components

Any custom flume components should inherit from the `org.apache.flume.instrumentation.MonitoredCounterGroup` class. The class should then provide getter methods for each of the metrics it exposes. See the code below. The `MonitoredCounterGroup` expects a list of attributes whose metrics are exposed by this class. As of now, this class only supports exposing metrics as long values.

```
public class SinkCounter extends MonitoredCounterGroup implements
    SinkCounterMBean {

    private static final String COUNTER_CONNECTION_CREATED =
        "sink.connection.creation.count";

    private static final String COUNTER_CONNECTION_CLOSED =
        "sink.connection.closed.count";

    private static final String COUNTER_CONNECTION_FAILED =
        "sink.connection.failed.count";

    private static final String COUNTER_BATCH_EMPTY =
        "sink.batch.empty";

    private static final String COUNTER_BATCH_UNDERFLOW =
        "sink.batch.underflow";

    private static final String COUNTER_BATCH_COMPLETE =
        "sink.batch.complete";

    private static final String COUNTER_EVENT_DRAIN_ATTEMPT =
        "sink.event.drain.attempt";

    private static final String COUNTER_EVENT_DRAIN_SUCCESS =
        "sink.event.drain.sucess";

    private static final String[] ATTRIBUTES = {
        COUNTER_CONNECTION_CREATED, COUNTER_CONNECTION_CLOSED,
        COUNTER_CONNECTION_FAILED, COUNTER_BATCH_EMPTY,
        COUNTER_BATCH_UNDERFLOW, COUNTER_BATCH_COMPLETE,
        COUNTER_EVENT_DRAIN_ATTEMPT, COUNTER_EVENT_DRAIN_SUCCESS
    };

    public SinkCounter(String name) {
        super(MonitoredCounterGroup.Type.SINK, name, ATTRIBUTES);
    }

    @Override
    public long getConnectionCreatedCount() {
        return get(COUNTER_CONNECTION_CREATED);
    }
}
```

```
public long incrementConnectionCreatedCount() {  
    return increment(COUNTER_CONNECTION_CREATED);  
}  
  
}
```

Tools

File Channel Integrity Tool

File Channel Integrity tool verifies the integrity of individual Events in the File channel and removes corrupted Events.

The tools can be run as follows:

```
$bin/flume-ng tool --conf ./conf FCINTEGRITYTOOL -l ./datadir
```

where datadir is the comma separated list of data directory to be verified.

Following are the options available

| Option Name | Description |
|-------------|---|
| h/help | Displays help |
| l/dataDirs | Comma-separated list of data directories which the tool must verify |

Event Validator Tool

Event validator tool can be used to validate the File Channel Event's in application specific way. The tool applies the user provider validation logic on each event and drop the event which do not confirm to the logic.

The tools can be run as follows:

```
$bin/flume-ng tool --conf ./conf FCINTEGRITYTOOL -l ./datadir -e  
org.apache.flume.MyEventValidator -DmaxSize 2000
```

where datadir is the comma separated list of data directory to be verified.

Following are the options available

| Option Name | Description |
|------------------|---|
| h/help | Displays help |
| l/dataDirs | Comma-separated list of data directories which the tool must verify |
| e/eventValidator | Fully Qualified Name of Event Validator Implementation. The jar must be on Flume classpath |

The Event validator implementation must implement EventValidator interface. It's recommended not to throw any exception from the implementation as they are treated as invalid events. Additional parameters can be passed to EventValidator implementation via -D options.

Let's see an example of simple size based Event Validator, which shall reject event's larger than maximum size specified.

```
public static class MyEventValidator implements EventValidator {  
  
    private int value = 0;  
  
    private MyEventValidator(int val) {  
        value = val;  
    }  
  
    @Override  
    public boolean validateEvent(Event event) {  
        return event.getBody().length() <= value;  
    }  
  
    public static class Builder implements EventValidator.Builder {  
  
        private int sizeValidator = 0;  
  
        @Override  
        public EventValidator build() {  
            return new DummyEventVerifier(sizeValidator);  
        }  
  
        @Override  
        public void configure(Context context) {  
            sizeValidator = context.getInteger("maxSize");  
        }  
    }  
}
```

Topology Design Considerations

Flume is very flexible and allows a large range of possible deployment scenarios. If you plan to use Flume in a large, production deployment, it is prudent to spend some time thinking about how to express your problem in terms of a Flume topology. This section covers a few considerations.

Is Flume a good fit for your problem?

If you need to ingest textual log data into Hadoop/HDFS then Flume is the right fit for your problem, full stop. For other use cases, here are some guidelines:

Flume is designed to transport and ingest regularly-generated event data over relatively stable, potentially complex topologies. The notion of “event data” is very broadly defined. To Flume, an event is just a generic blob of bytes. There are some limitations on how large an event can be - for instance, it cannot be larger than what you can store in memory or on disk on a single machine - but in practice, flume events can be everything from textual log entries to image files. The key property of an event is that they are generated in a continuous, streaming fashion. If your data is not regularly generated (i.e. you are trying to do a single bulk load of data into a Hadoop cluster) then Flume will still work, but it is probably overkill for your situation. Flume likes relatively stable topologies. Your topologies do not need to be immutable, because Flume can deal with changes in topology without losing data and can also tolerate periodic reconfiguration due to fail-over or provisioning. It probably won't work well if you plan to change topologies every day, because reconfiguration takes some thought and overhead.

Flow reliability in Flume

The reliability of a Flume flow depends on several factors. By adjusting these factors, you can achieve a wide array of reliability options with Flume.

What type of channel you use. Flume has both durable channels (those which will persist data to disk) and non durable channels (those which will lose data if a machine fails). Durable channels use disk-based storage, and data stored in such channels will persist across machine restarts or non disk-related failures.

Whether your channels are sufficiently provisioned for the workload. Channels in Flume act as buffers at various hops. These buffers have a fixed capacity, and once that capacity is full you will create back pressure on earlier points in the flow. If this pressure propagates to the source of the flow, Flume will become unavailable and may lose data.

Whether you use redundant topologies. Flume lets you replicate flows across redundant topologies. This can provide a very easy source of fault tolerance and one which overcomes both disk or machine failures.

The best way to think about reliability in a Flume topology is to consider various failure scenarios and their outcomes. What happens if a disk fails? What happens if a machine fails? What happens if your terminal sink (e.g. HDFS) goes down for some time and you have back pressure? The space of possible designs is huge, but the underlying questions you need to ask are just a handful.

Flume topology design

The first step in designing a Flume topology is to enumerate all sources and destinations (terminal sinks) for your data. These will define the edge points of your topology. The next consideration is whether to introduce intermediate aggregation tiers or event routing. If you are collecting data from a large number of sources, it can be helpful to aggregate the data in order to simplify ingestion at the terminal sink. An aggregation tier can also smooth out burstiness from sources or unavailability at sinks, by acting as a buffer. If you are routing data between different locations, you may also want to split flows at various points: this creates sub-topologies which may themselves include aggregation points.

Sizing a Flume deployment

Once you have an idea of what your topology will look like, the next question is how much hardware and networking capacity is needed. This starts by quantifying how much data you generate. That is not always a simple task! Most data streams are bursty (for instance, due to diurnal patterns) and potentially unpredictable. A good starting point is to think about the maximum throughput you'll have in each tier of the topology, both in terms of *events per second* and *bytes per second*. Once you know the required throughput of a given tier, you can calculate a lower bound on how many nodes you require for that tier. To determine attainable throughput, it's best to experiment with Flume on your hardware, using synthetic or sampled event

data. In general, disk-based channels should get 10's of MB/s and memory based channels should get 100's of MB/s or more. Performance will vary widely, however depending on hardware and operating environment.

Sizing aggregate throughput gives you a lower bound on the number of nodes you will need to each tier. There are several reasons to have additional nodes, such as increased redundancy and better ability to absorb bursts in load.

Troubleshooting

Handling agent failures

If the Flume agent goes down, then the all the flows hosted on that agent are aborted. Once the agent is restarted, then flow will resume. The flow using file channel or other stable channel will resume processing events where it left off. If the agent can't be restarted on the same hardware, then there is an option to migrate the database to another hardware and setup a new Flume agent that can resume processing the events saved in the db. The database HA futures can be leveraged to move the Flume agent to another host.